

Active Learning

Recap

ThermalWear

Clothing Style (Type + Color)

+

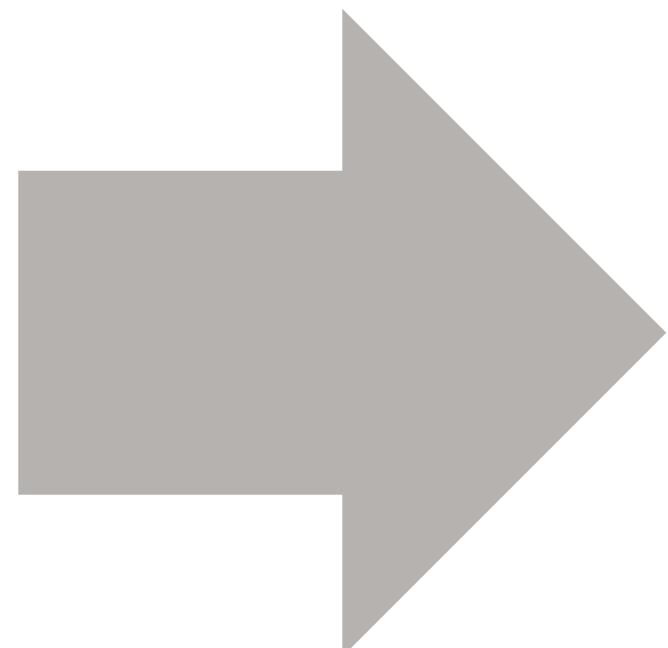
Local Temperature

+

Local Humidity

=

People's body temperature



ComfyWear

Clothing Style (Type only)

+

Local Temperature

+

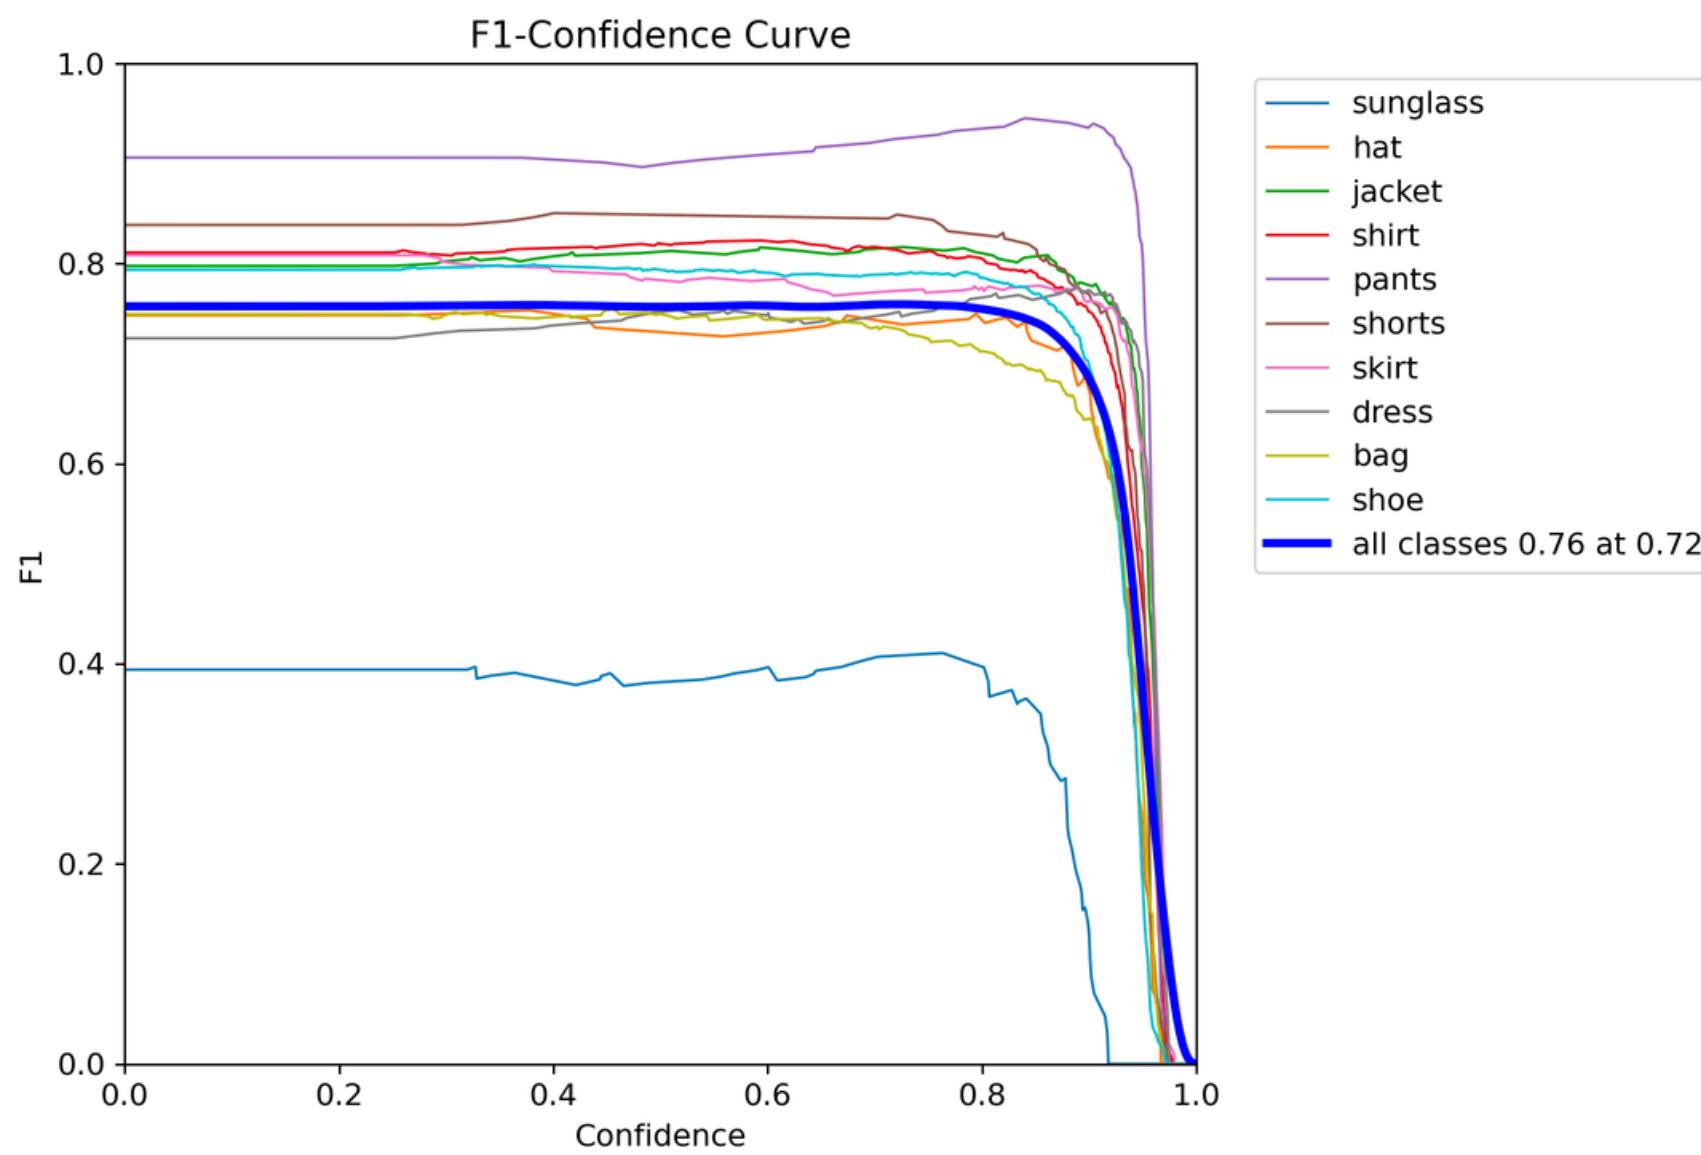
Local Humidity

=

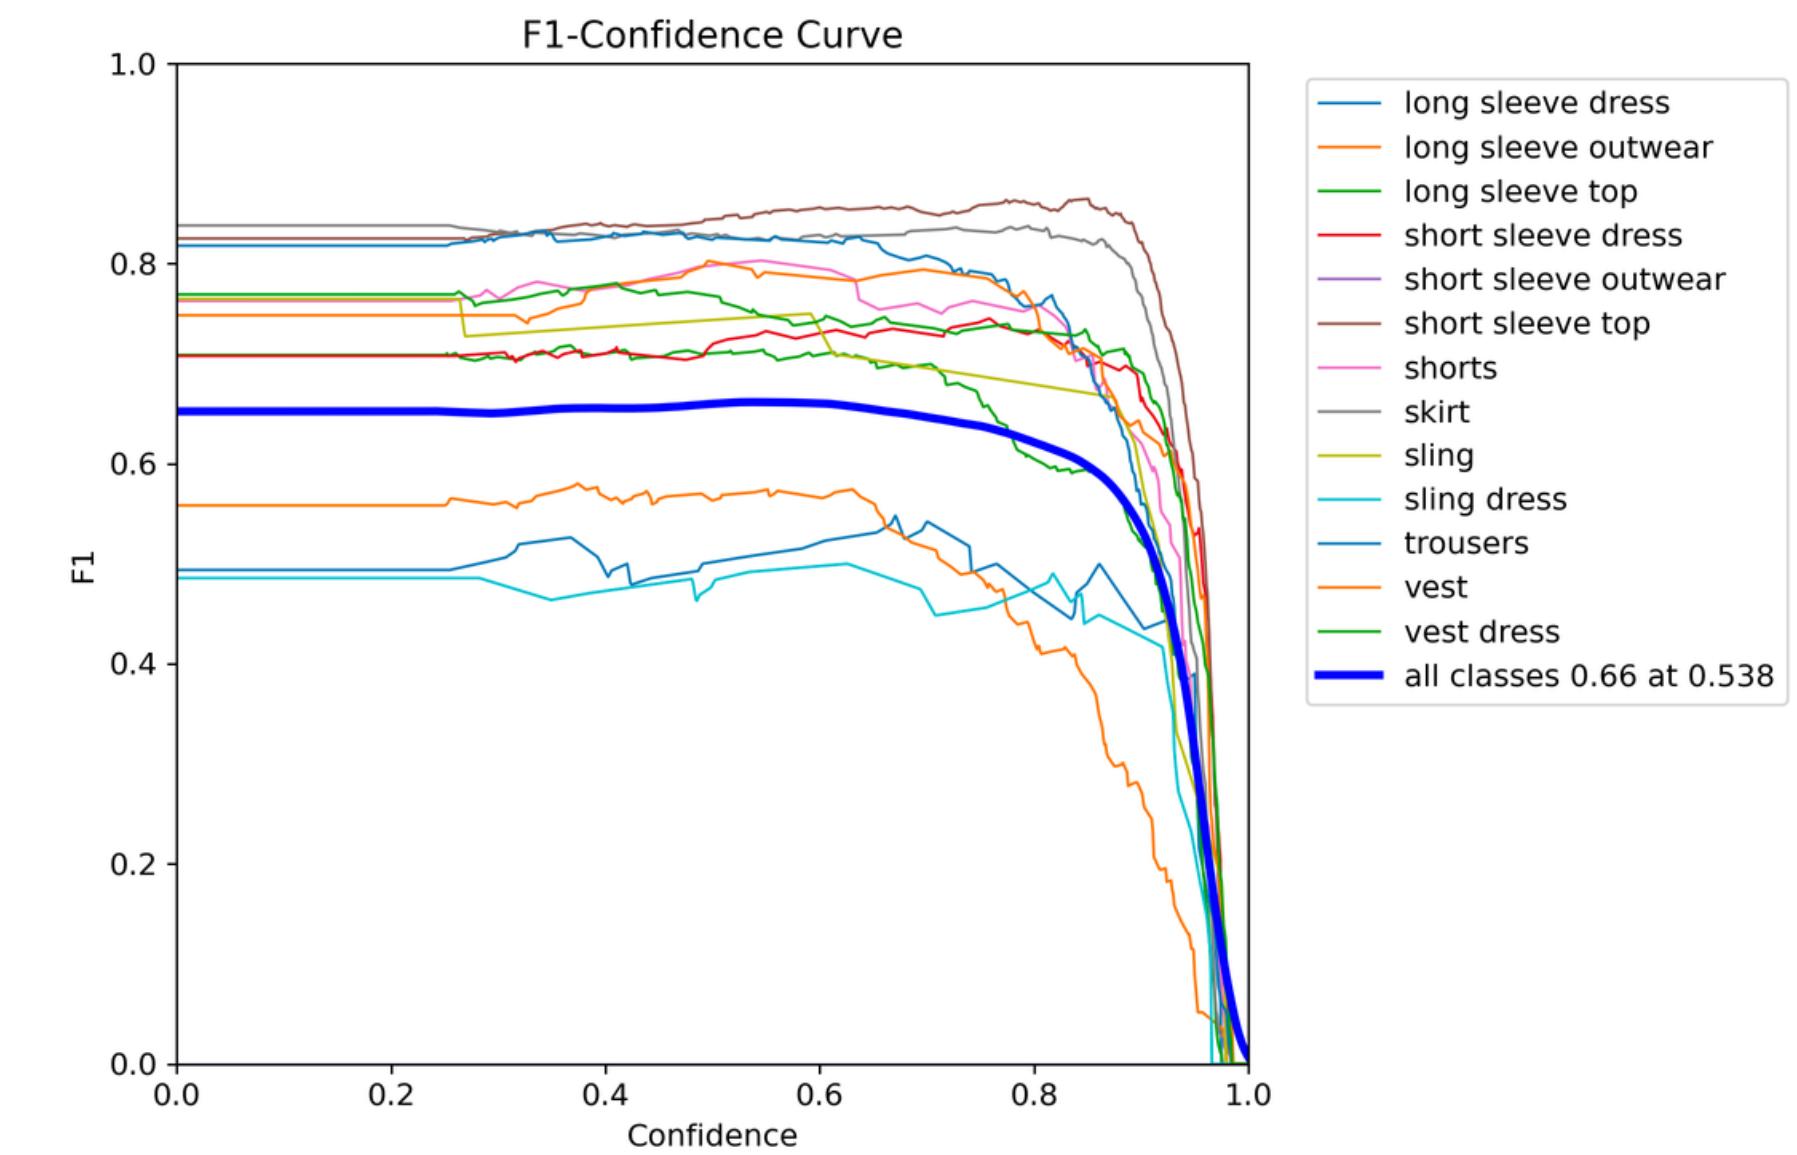
Comfort

F1-Confidence Curve

Colorful



Deepfashion2



Soft-Biometric (UPAR)

20 epochs-MultiLabels

BCEWithLogitsLoss():
Sigmoid + BCELoss

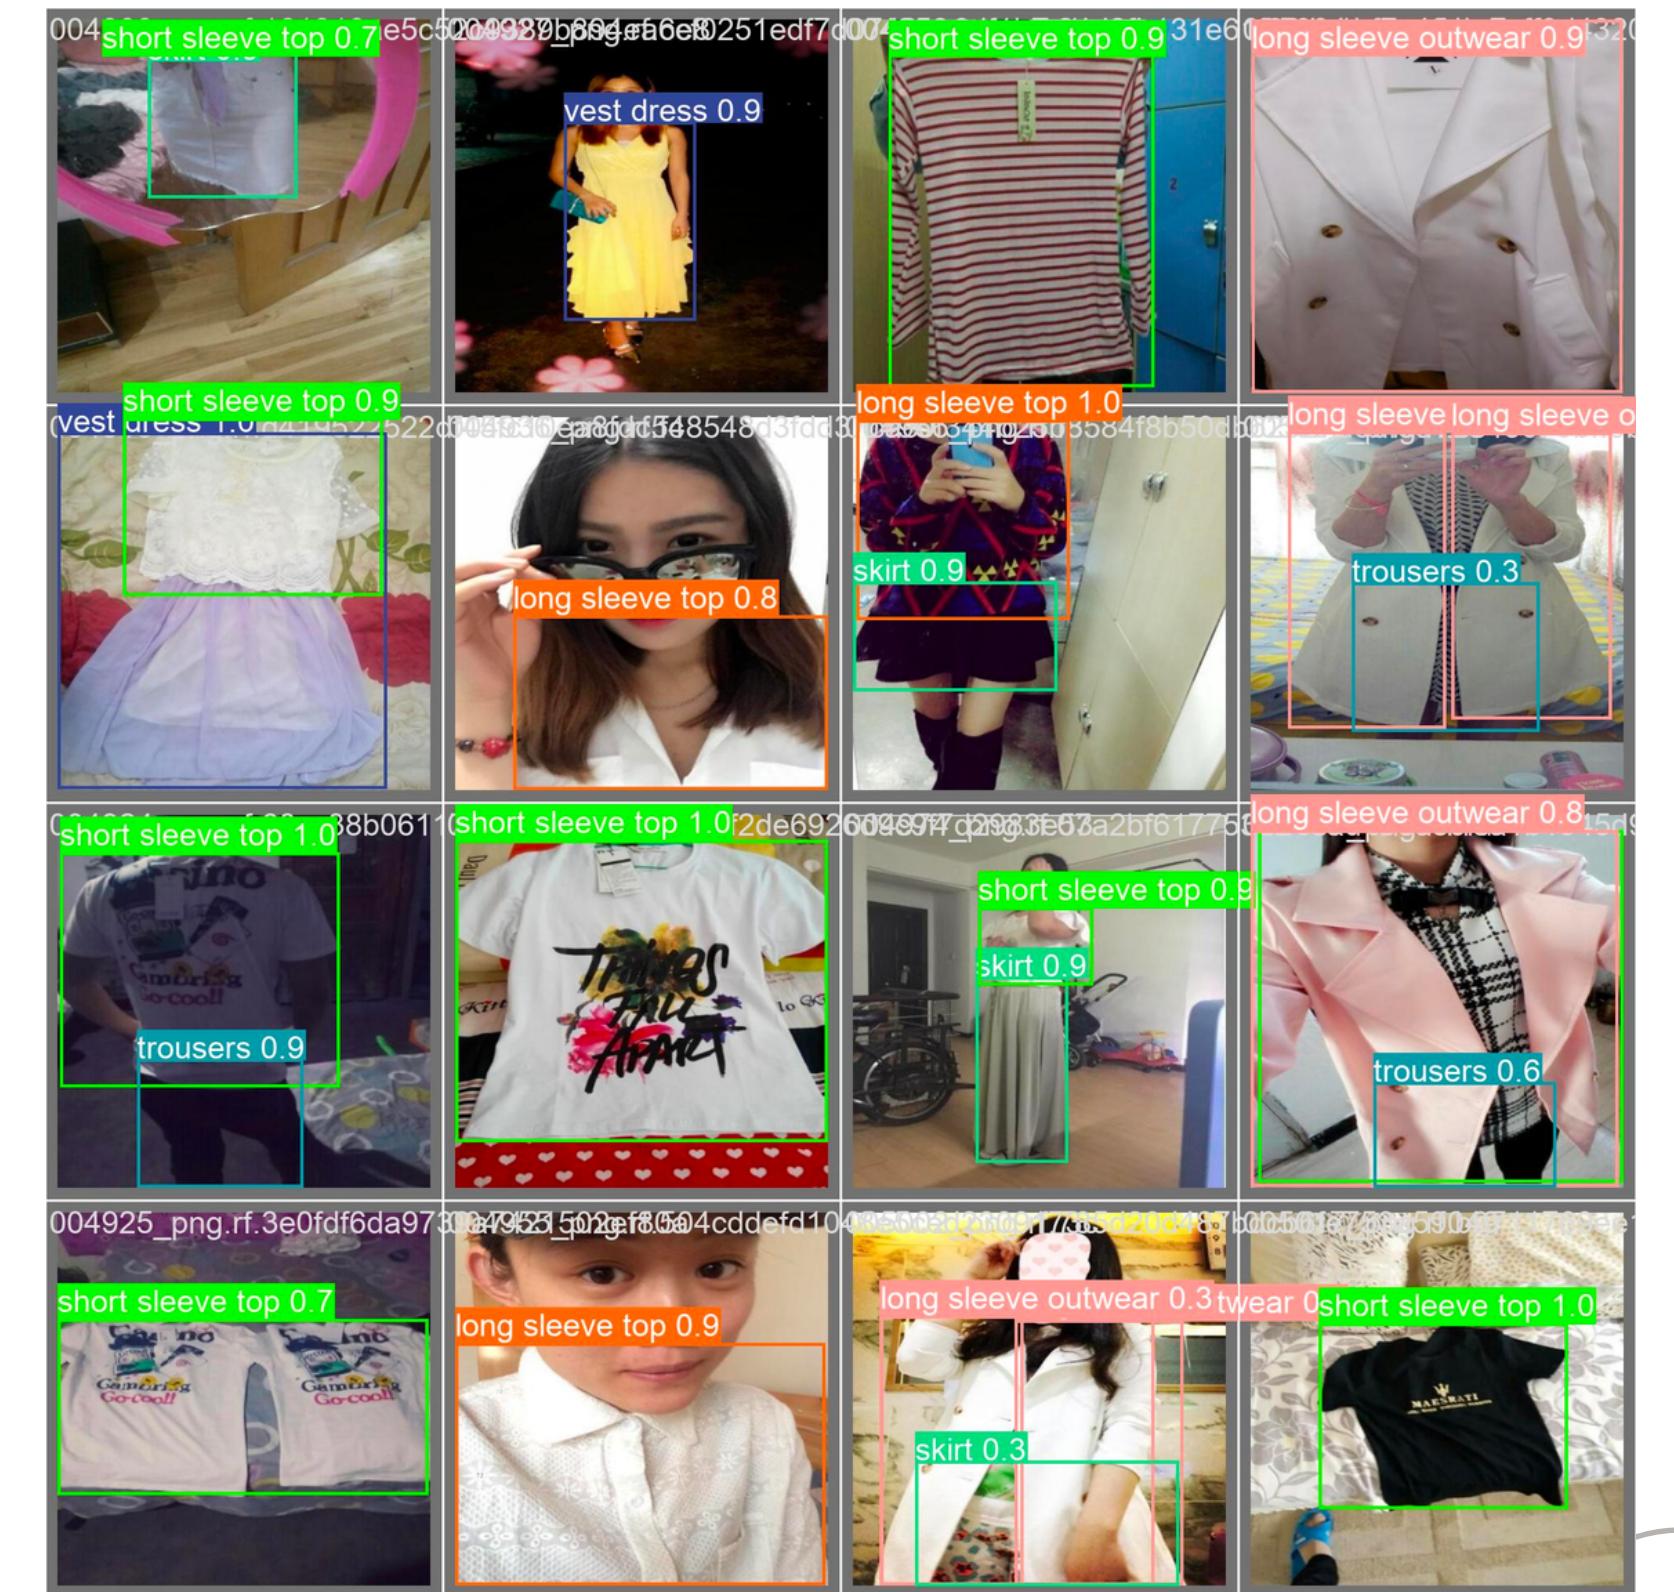


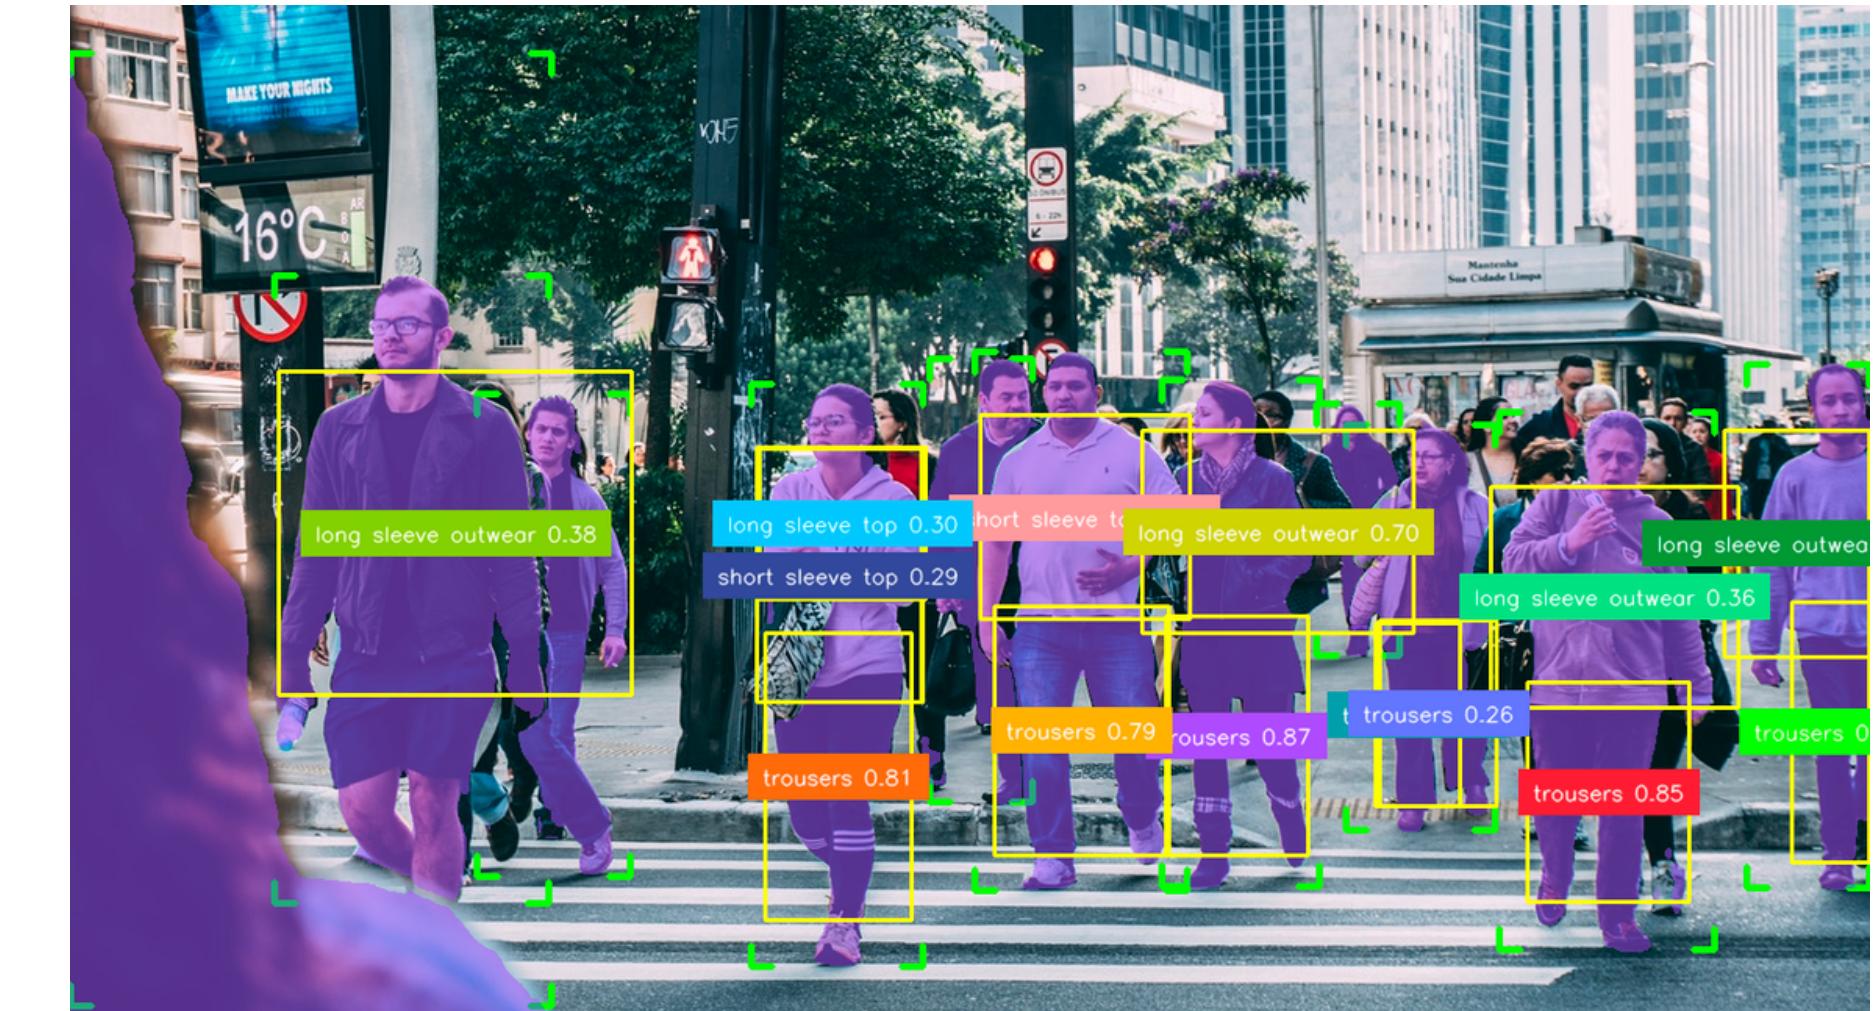
Colorful



Example

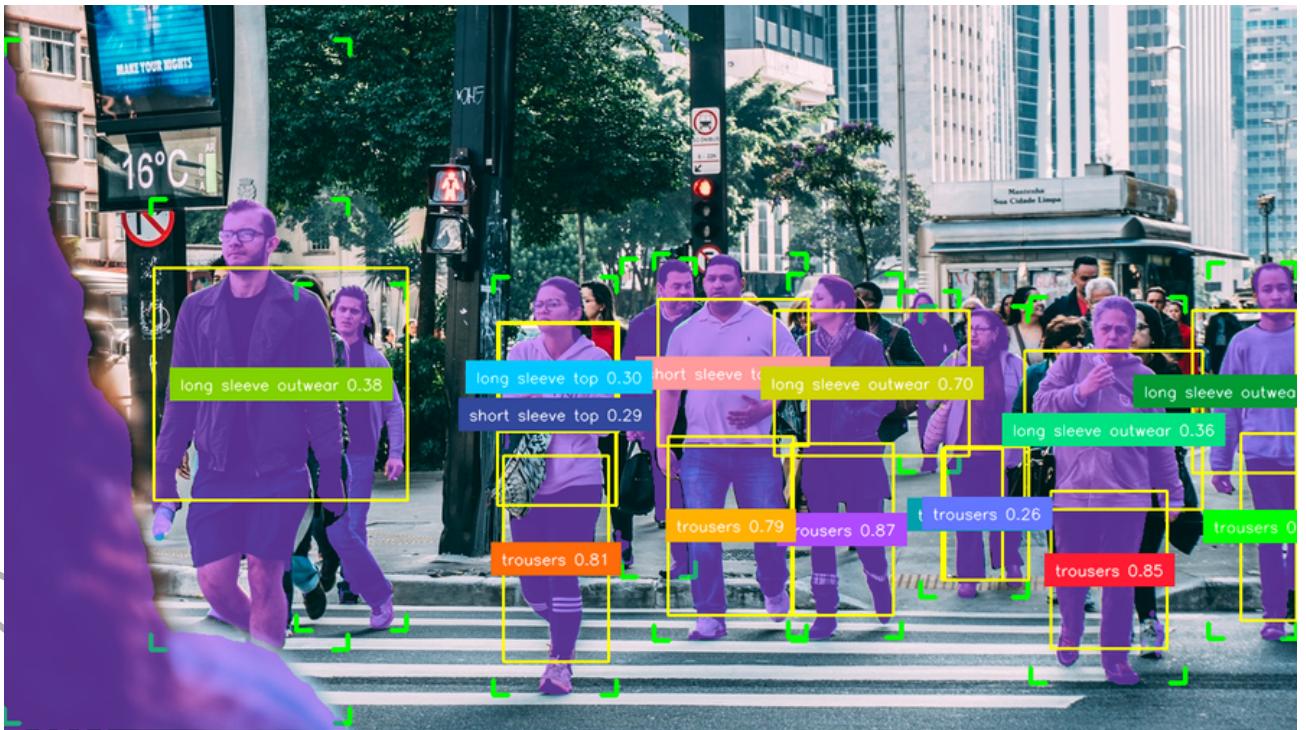
Deepfashion2





Number of Prediction

14



24



Feedback

- **Collect 1-5 comfort level**
- **Detect the color of clothes**
- **People Re-Identification -> Create own dataset**
- **Human reinforcement**



The background features a minimalist design with light gray, thin-lined abstract shapes. In the upper right quadrant, there is a series of concentric circles that decrease in size towards the center. The lower left quadrant contains several wavy, organic lines that curve upwards and outwards. The overall aesthetic is clean and modern, using negative space and simple lines to create a sense of depth and movement.

Today's topic

Human Reinforcement

Reinforcement Learning

A machine learning paradigm that aims to train an agent (a software program or system) to make optimal decisions or take optimal actions in an environment to maximize a reward signal or cumulative reward over time

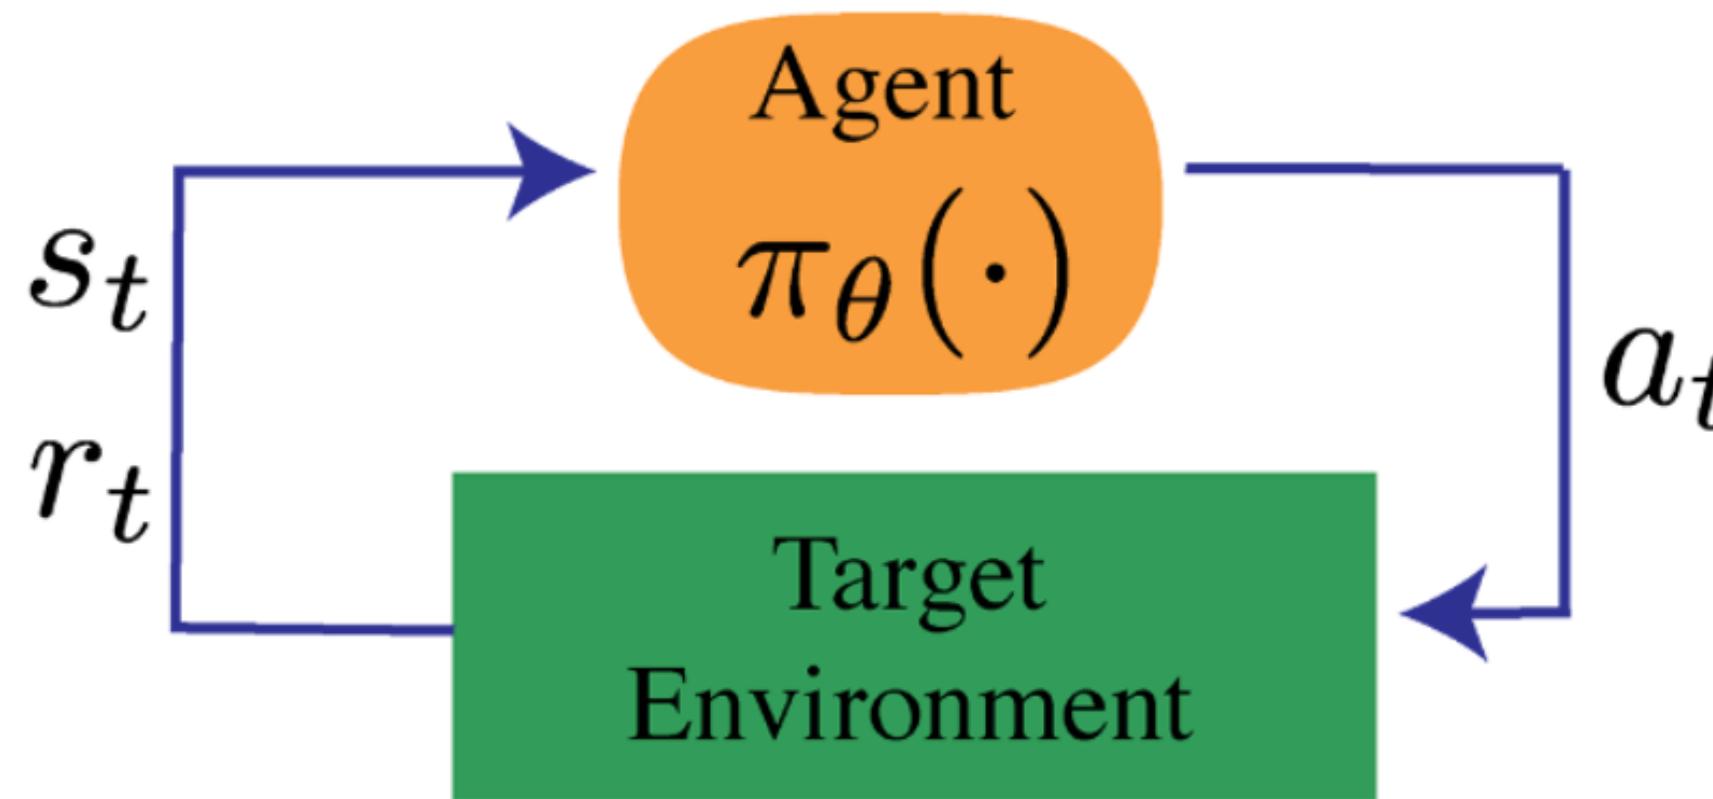
Key Elements

- **Agent:** The Learner
- **Environment:** The world or system in which the agent operates and takes actions.
- **State:** The current situation or condition of the environment that the agent perceives.

Key Elements

- **Action:** The choices or decisions that the agent can make to change its state in the environment.
- **Reward:** The feedback signal (positive or negative)
- **Policy:** A strategy or set of rules that an agent follows to make decisions.

How does RL work?



Some notation:

s_t : state

r_t : reward

a_t : action

$a_t \sim \pi_\theta(s_t)$: policy

RL vs SL?

Similarities

- Goal
- Learning from data
- Optimization

Differences

- Input/Output Mapping
 - State -> Action, $X \rightarrow Y$
- Feedback Mechanism
- Exploration vs. Exploitation

Problems

How can we create a loss function for:

- What is funny?
- What is *ethical*?
- What is *safe*?
- What is taste?

Reinforcement Learning Human Feedback (RLHF)

RLHF

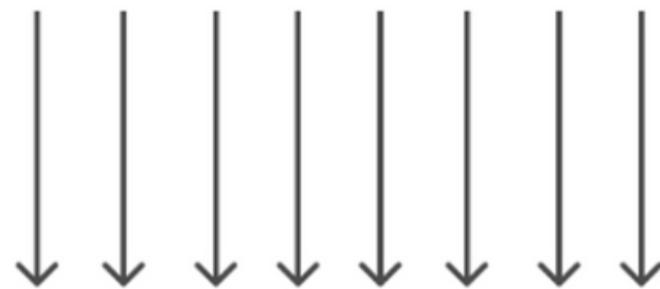
The process of incorporating human feedback into machine learning algorithms to enhance learning efficiency and accuracy.

Mostly used in NLP tasks.

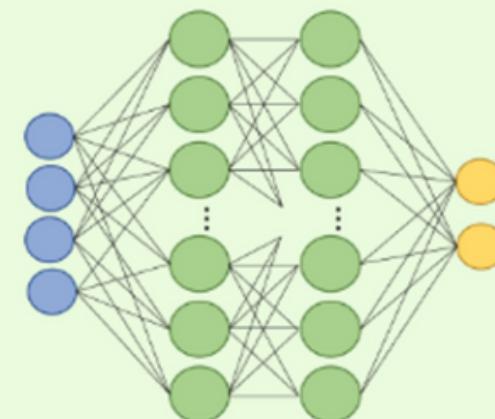
Prompts Dataset



Sample many prompts



Initial Language Model



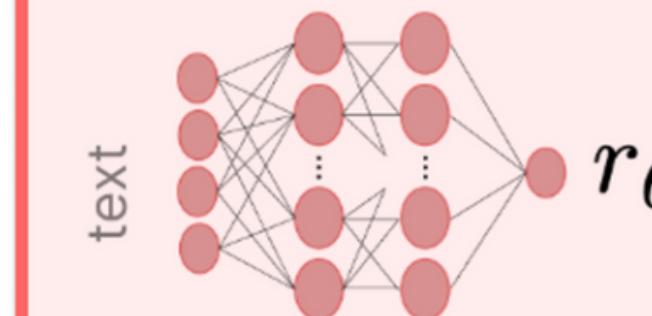
Generated text
Lorem ipsum dolor
sit amet, consectetur
adipiscing elit. Aen
Donec quam felis
vulputate eget, arc
Nam quam nunc
eros faucibus tincid
luctus pulvinar, her



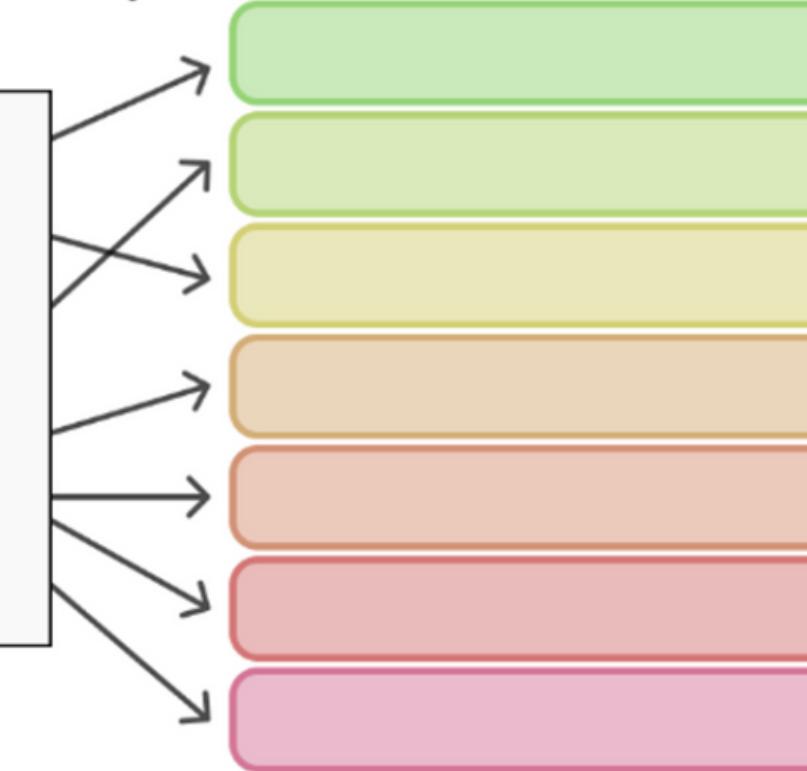
Generated text

Train on
{sample, reward} pairs

Reward (Preference) Model

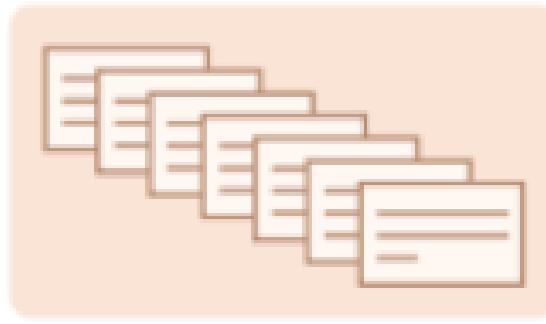


Outputs are ranked
(relative, ELO, etc.)

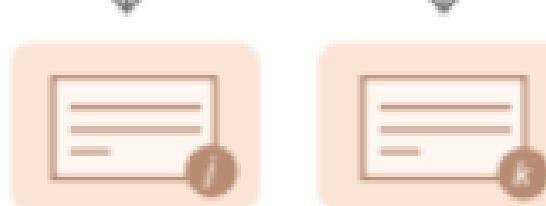


① Collect human feedback

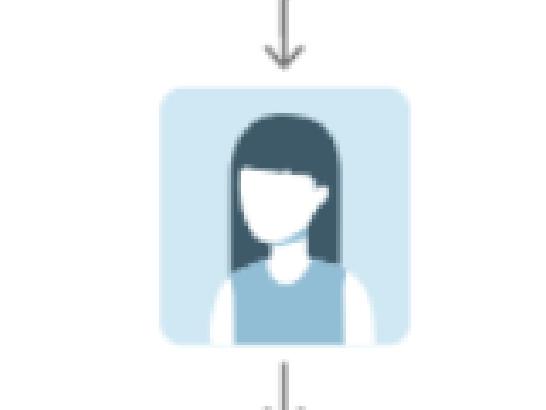
A Reddit post is sampled from the Reddit TL;DR dataset.



Various policies are used to sample a set of summaries.



Two summaries are selected for evaluation.



A human judges which is a better summary of the post.

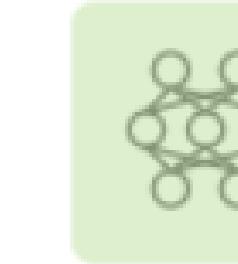
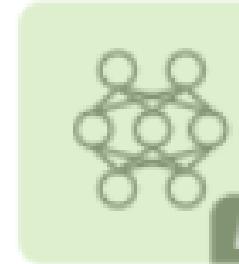
"j is better than k"

② Train reward model

One post with two summaries judged by a human are fed to the reward model.



The reward model calculates a reward r for each summary.



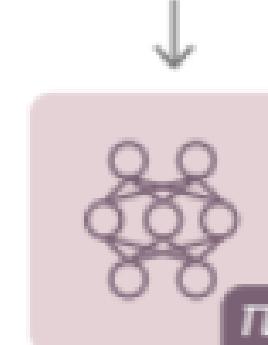
The loss is calculated based on the rewards and human label, and is used to update the reward model.

$$\text{loss} = \log(a(r_j - r_k))$$

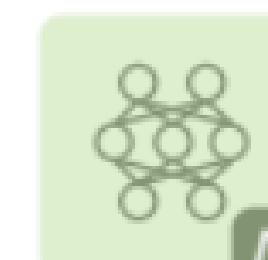
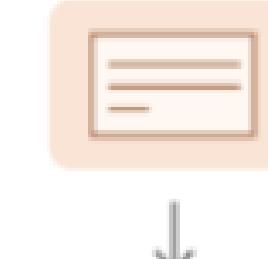
"j is better than k"

③ Train policy with PPO

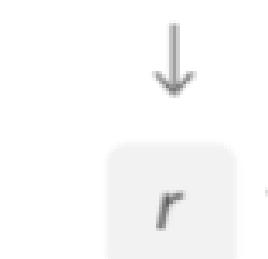
A new post is sampled from the dataset.



The policy π generates a summary for the post.



The reward model calculates a reward for the summary.



The reward is used to update the policy via PPO.

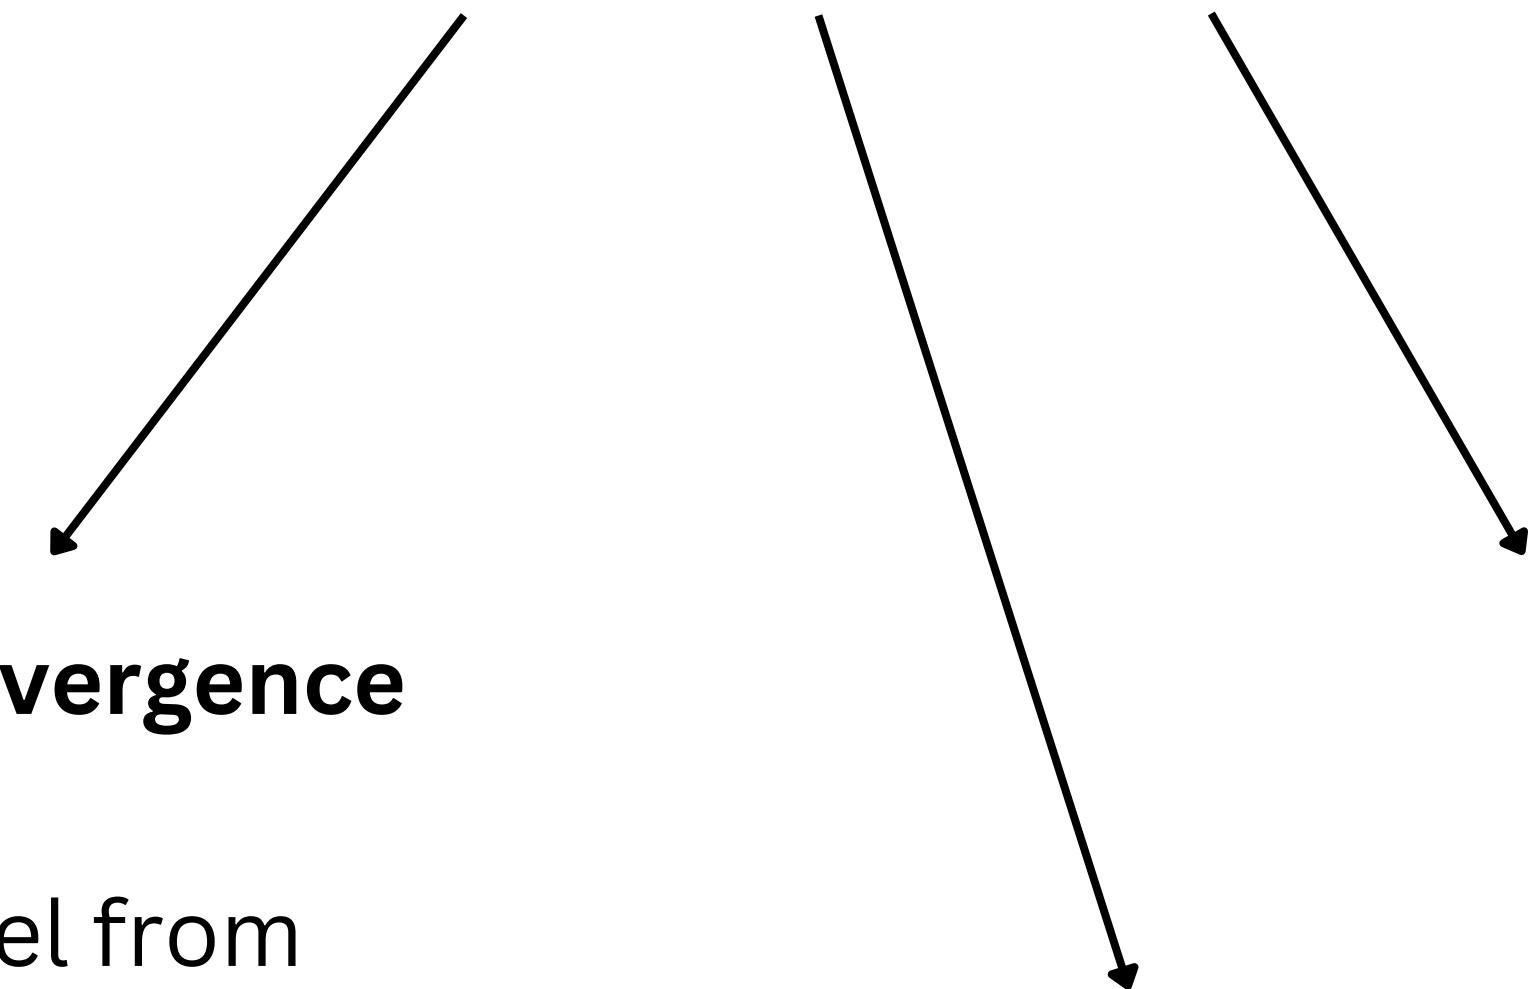
Fine-Tuning RL

Kullback-Leibler Divergence

Prevent the model from
deviating too much from the
original pre-trained model.

PPO

Reward model





So should we use the RL?

CNN is better

So should we use the RL?



Reward Shaping Complexity

CNN is better

So should we use the RL?

Reward Shaping Complexity

CNN is better

Lack of Sequential Decision Making

So should we use the RL?

Reward Shaping Complexity

CNN is better

Lack of Sequential Decision Making

So should we use the RL?

Pretrained did well already

How to train LLM?

How to train LLM?

1. **Pre-training on Text Data:** Using a self-supervised objective like predicting masked words. (**Unsupervised**)
2. **Fine-tuning on Labeled Data:** The pre-trained LLM is then further trained on labeled datasets for specific tasks (**Supervised**)
3. **Getting Human Feedback:** Step 1 in RLHF
4. **Reinforcement Learning:** Step 2 in RLHF
5. **Creating New Datasets:** Use human to make high-quality output references
6. **Dataset Fine-tuning:** Train on dataset from 5
7. Repeat 3-6

(Google's LaMDA, Anthropic's Constitutional AI, and OpenAI's InstructGPT)

Extract The Color From detection

KMeans Clustering

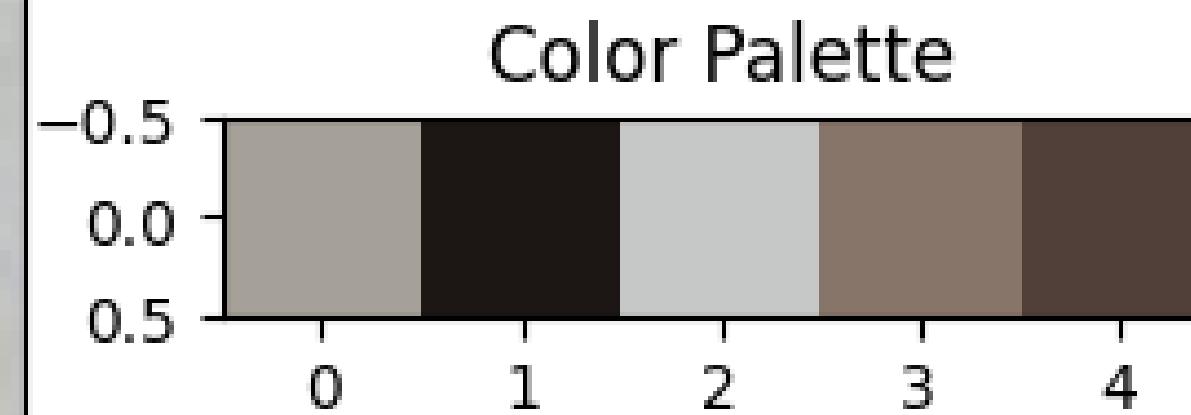
```
from sklearn.cluster import KMeans
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.image as mpimg

image = mpimg.imread('person/person/im.jpg4.jpg')
image = image / 255.0
w, h, d = tuple(image.shape)
pixel = np.reshape(image, (w * h, d))

n_colors = 5
model = KMeans(n_clusters=n_colors, random_state=42)
labels = model.fit_predict(pixel)
colors = model.cluster_centers_
colors = (colors * 255).astype(np.uint8)

plt.subplot(1, 2, 1)
plt.imshow(image)
plt.title('Original Image')
```

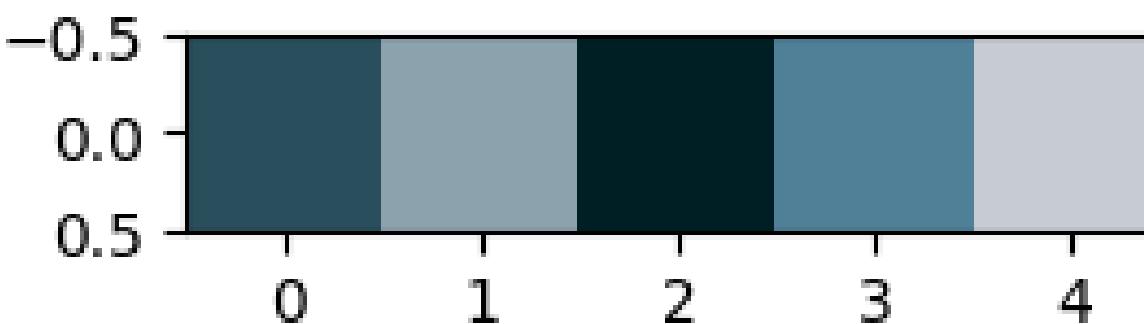
- Normalize the image
- Flatten the pixel data into 2D
- Initialize k centroids (**number of colors**)
- Assign data points to clusters
- Update centroids
- Repeat



Original Image



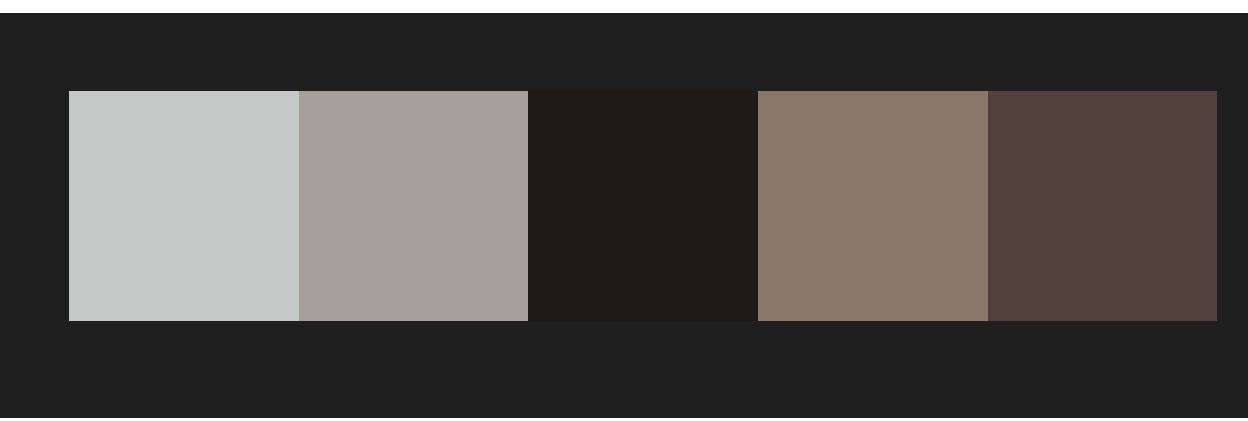
Color Palette



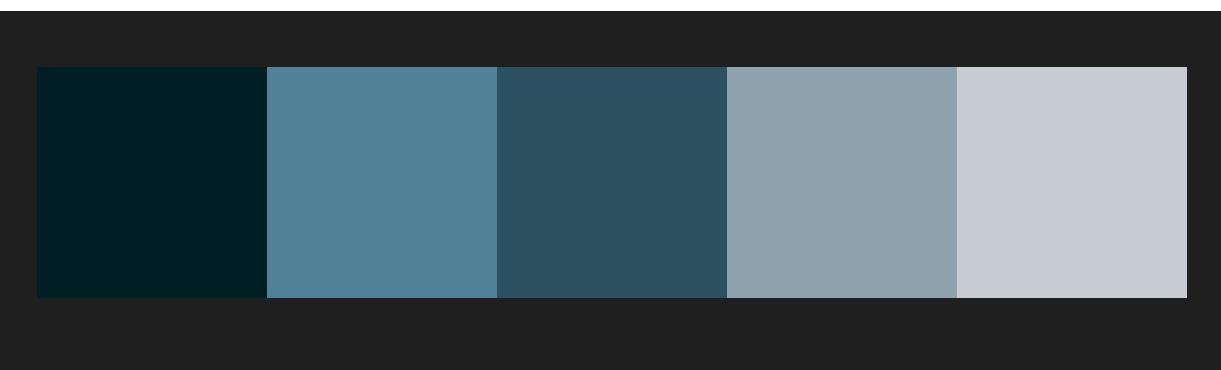
Pylette

```
from Pylette import extract_colors  
palette = extract_colors(image='person/person/im.jpg4.jpg', palette_size=5, resize=True)
```

- Pylette is optimized for color extraction and palette generation, making it more specialized and potentially more accurate for this task compared to the general-purpose k-means algorithm.
- Uses a perceptual color clustering algorithm based on the CIEDE2000 color difference formula

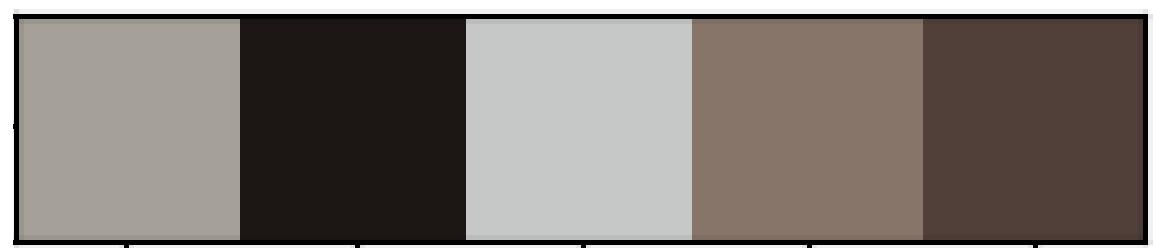


Original Image



Compare the result

KMeans



Pylette



Compare the result

KMeans



Pylette



Original Image



Problem

Color Mapping

Our Label

- (255, 0, 0) -> Red
- (255, 165, 0) -> Orange
- (255, 255, 0) -> Yellow
- (0, 128, 0) -> Green
- (0, 0, 255) -> Blue
- (75, 0, 130) -> indigo
- (238, 130, 238) -> Violet
- White, Black

What we extracted

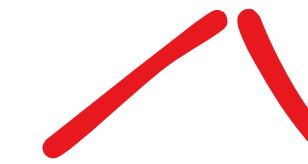
- (255, 102, 102)
- (255, 204, 102)
- (255, 255, 102)
- (102, 204, 102)
- Doesn't matched any label

So, how can we solve that?

Map the RGB to our closest label

Manually map by distance

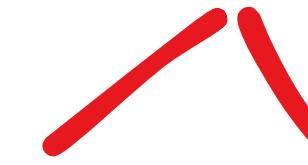
rgb



```
def euclidean_distance(color1, color2):  
    return sqrt(sum((c1 - c2) ** 2 for c1, c2 in zip(color1, color2)))
```

Manually map by distance

rgb

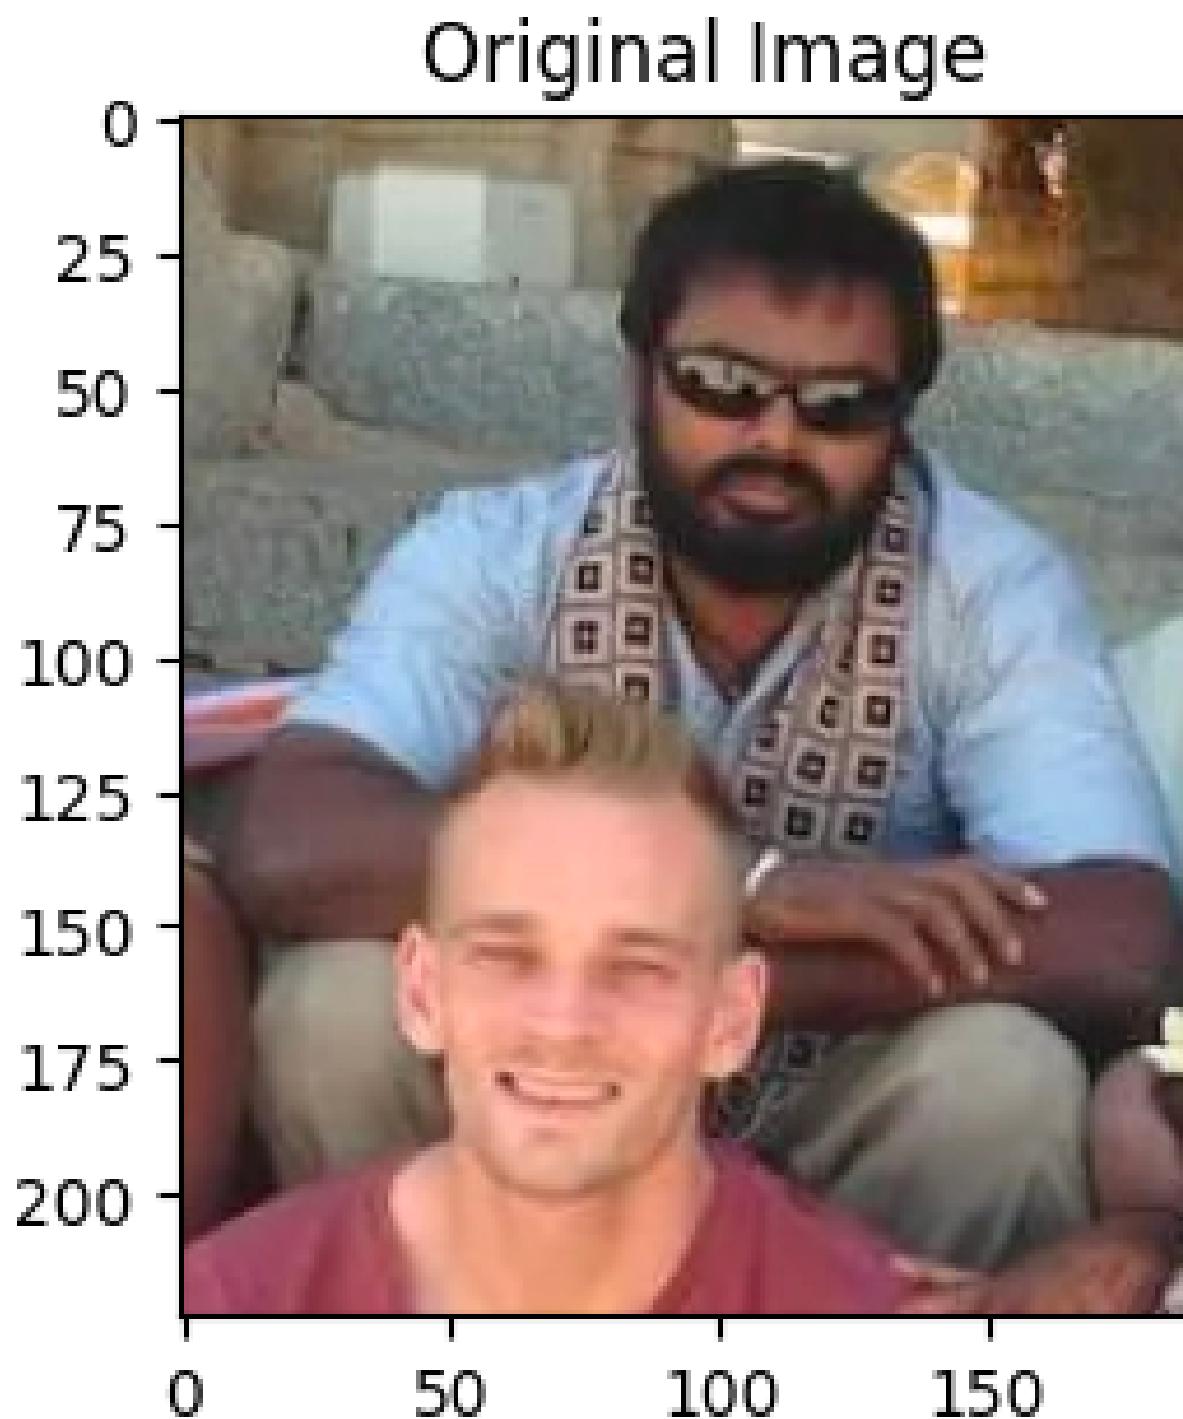


```
def euclidean_distance(color1, color2):  
    return sqrt(sum((c1 - c2) ** 2 for c1, c2 in zip(color1, color2)))
```

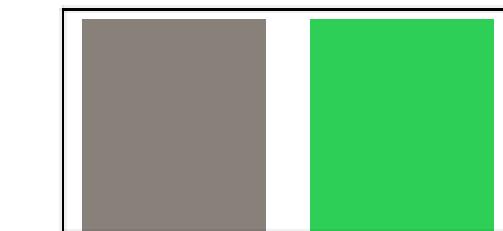
Might not really accurate

```
def map_to_rainbow_color(rgb):
    rainbow_colors = {
        'red': (255, 0, 0),
        'light_red': (255, 102, 102),
        'dark_red': (153, 0, 0),
        'orange': (255, 165, 0),
        'light_orange': (255, 204, 102),
        'dark_orange': (204, 102, 0),
        'yellow': (255, 255, 0),
        'light_yellow': (255, 255, 102),
        'dark_yellow': (204, 204, 0),
        'green': (0, 128, 0),
        'light_green': (102, 204, 102),
        'dark_green': (0, 102, 0),
        'blue': (0, 0, 255),
        'light_blue': (102, 102, 255),
        'dark_blue': (0, 0, 153),
        'indigo': (75, 0, 130),
        'light_indigo': (138, 43, 226),
        'dark_indigo': (49, 0, 98),
        'violet': (238, 130, 238),
        'light_violet': (255, 187, 255),
        'dark_violet': (147, 112, 219),
        'black': (0, 0, 0),
        'white': (255, 255, 255)
    }
    closest_color = min(rainbow_colors, key=lambda color: euclidean_distance(rgb, rainbow_colors[color]))
    return closest_color, rainbow_colors[closest_color]
```

Result

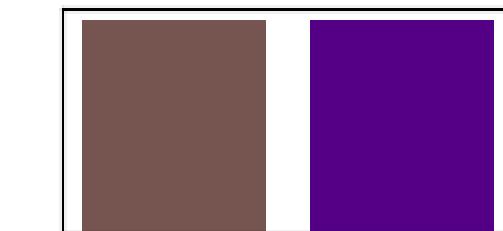


RGB: (135, 129, 123) -> Mapped Color: light_green



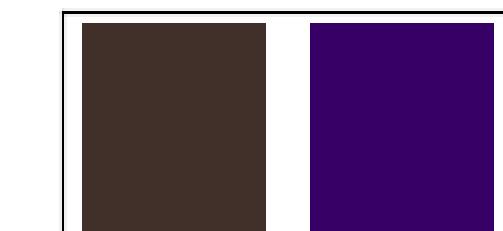
Original Mapped Color

RGB: (113, 85, 81) -> Mapped Color: indigo



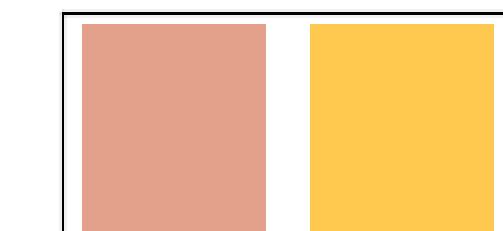
Original Mapped Color

RGB: (62, 48, 43) -> Mapped Color: dark_indigo



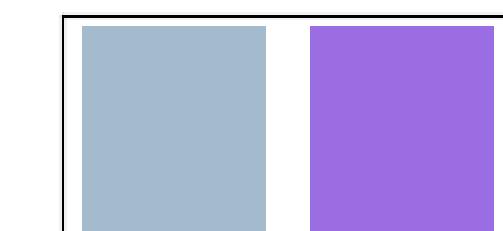
Original Mapped Color

RGB: (217, 163, 142) -> Mapped Color: light_orange



Original Mapped Color

RGB: (168, 186, 203) -> Mapped Color: dark_violet



Original Mapped Color

The result isn't good.

It might be because of a lack of mapping rgb.

Webcolors

Webcolors

```
from webcolors import (
    CSS3_HEX_TO_NAMES,
    hex_to_rgb,
)
```



138 colors

Webcolors

```
from scipy.spatial import KDTree
```

```
def convert_rgb_to_names(rgb_tuple):
    css3_db = css3_hex_to_names
    names = []
    rgb_values = []
    for color_hex, color_name in css3_db.items():
        names.append(color_name)
        rgb_values.append(hex_to_rgb(color_hex))

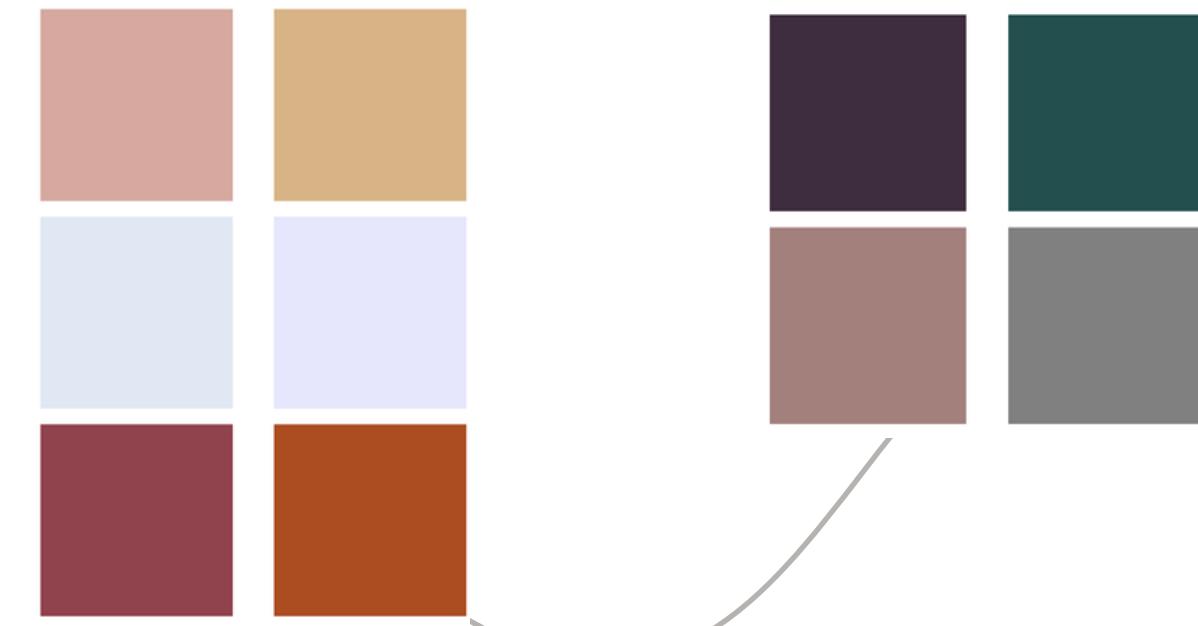
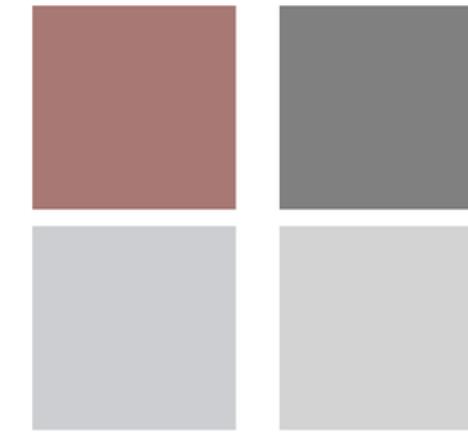
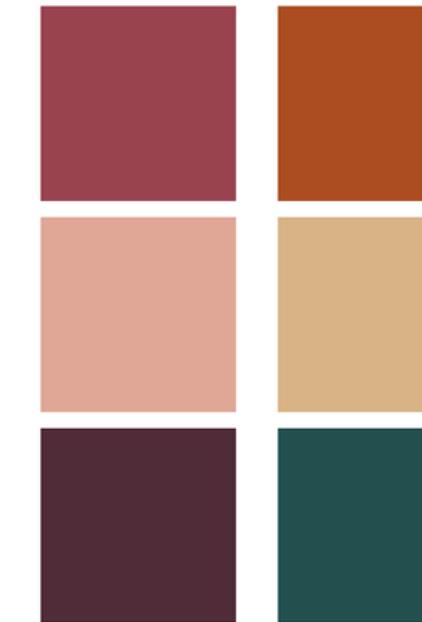
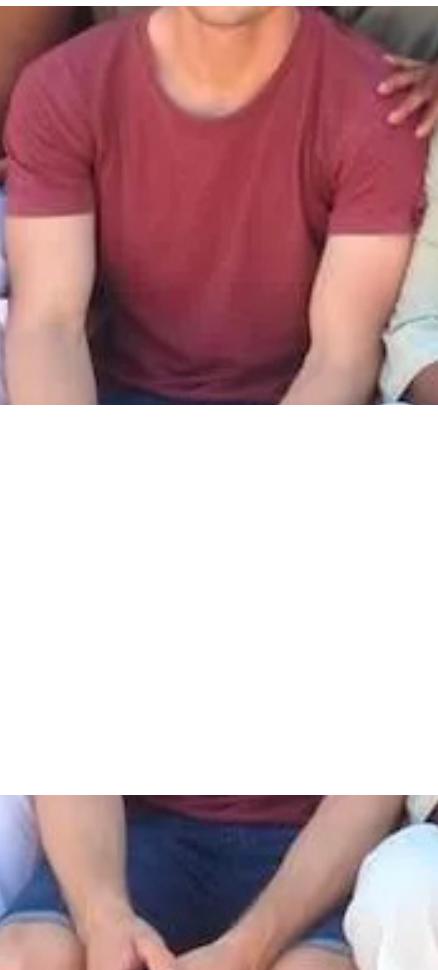
    kdt_db = KDTree(rgb_values)
    distance, index = kdt_db.query(rgb_tuple)
    return f'closest match: {names[index]}
```

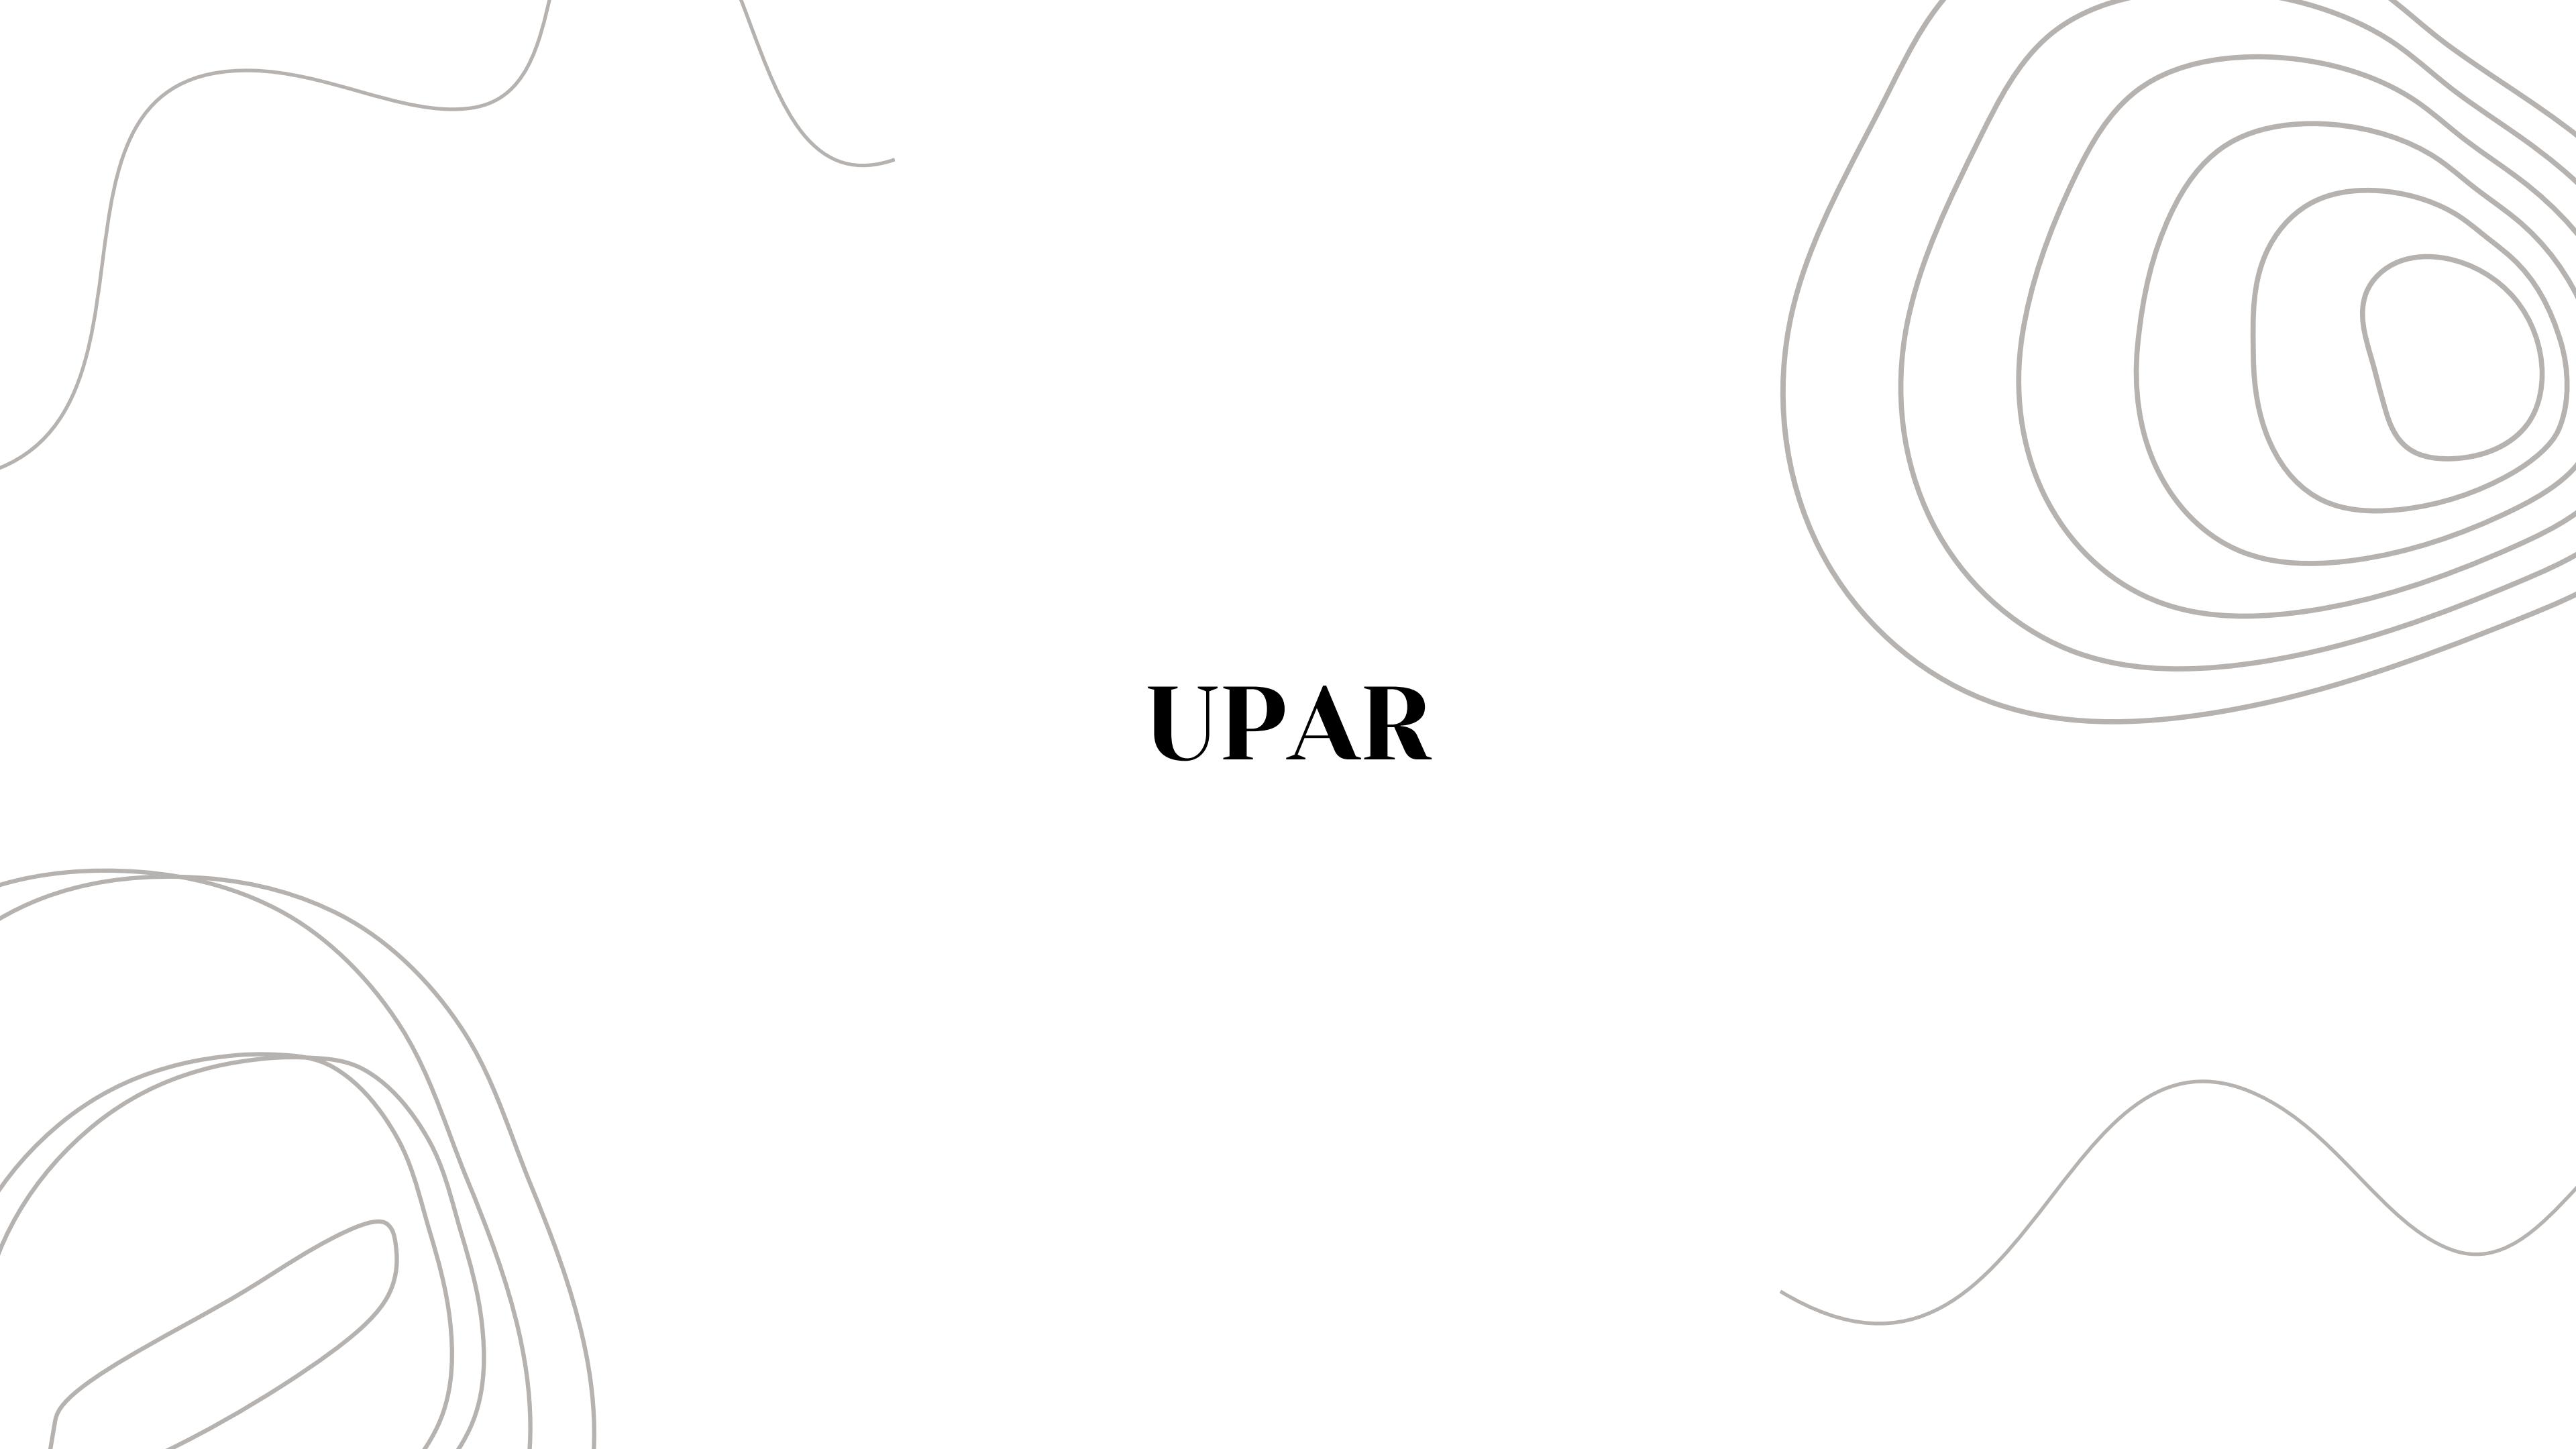
Performs a nearest neighbor search in the K-D Tree to find the closest RGB value to the given `rgb_tuple`.

Using Euclidean distance

RESULTS

RESULTS





The background features a minimalist design with abstract, light-grey line art. It includes several sets of concentric circles of varying sizes, some with irregular edges, and some organic, flowing lines that resemble stylized leaves or waves.

UPAR

Compare Models

Resnet50

- Same model as first place in the UPAR Challenge Competition.
- Trained 50 epochs
- ~37% accuracy

MobileNet

- Trained 50 epochs
- ~95% accuracy
- Haven't evaluate yet

Person Re-Identification

A Comprehensive Overview of Person Re-Identification Approaches (2020)

Steps

1. **Local information extraction:** Extract feature from different body blocks/parts.
2. **Global information extraction:** Extract feature from multiple non-overlapped cameras.
3. **Metric Learning:** Find an optimized distance metric, designing new distance functions and loss functions.

Steps

4. Post processing: Applying re-ranking algorithms on the initial ranking list obtained from the re-identification model.



Steps

4. **Efficiency improvement:** All about the models.
5. **Labeling cost reduction:** Data augmentation, unsupervised domain adaptation, semi-supervised learning. (reduce the amount of labeled training data required).
6. **Data type extension:** Images -> videos, depth images, infrared images, low-resolution images.



A Comprehensive Overview of Person Re-Identification Approaches (2020)

Steps

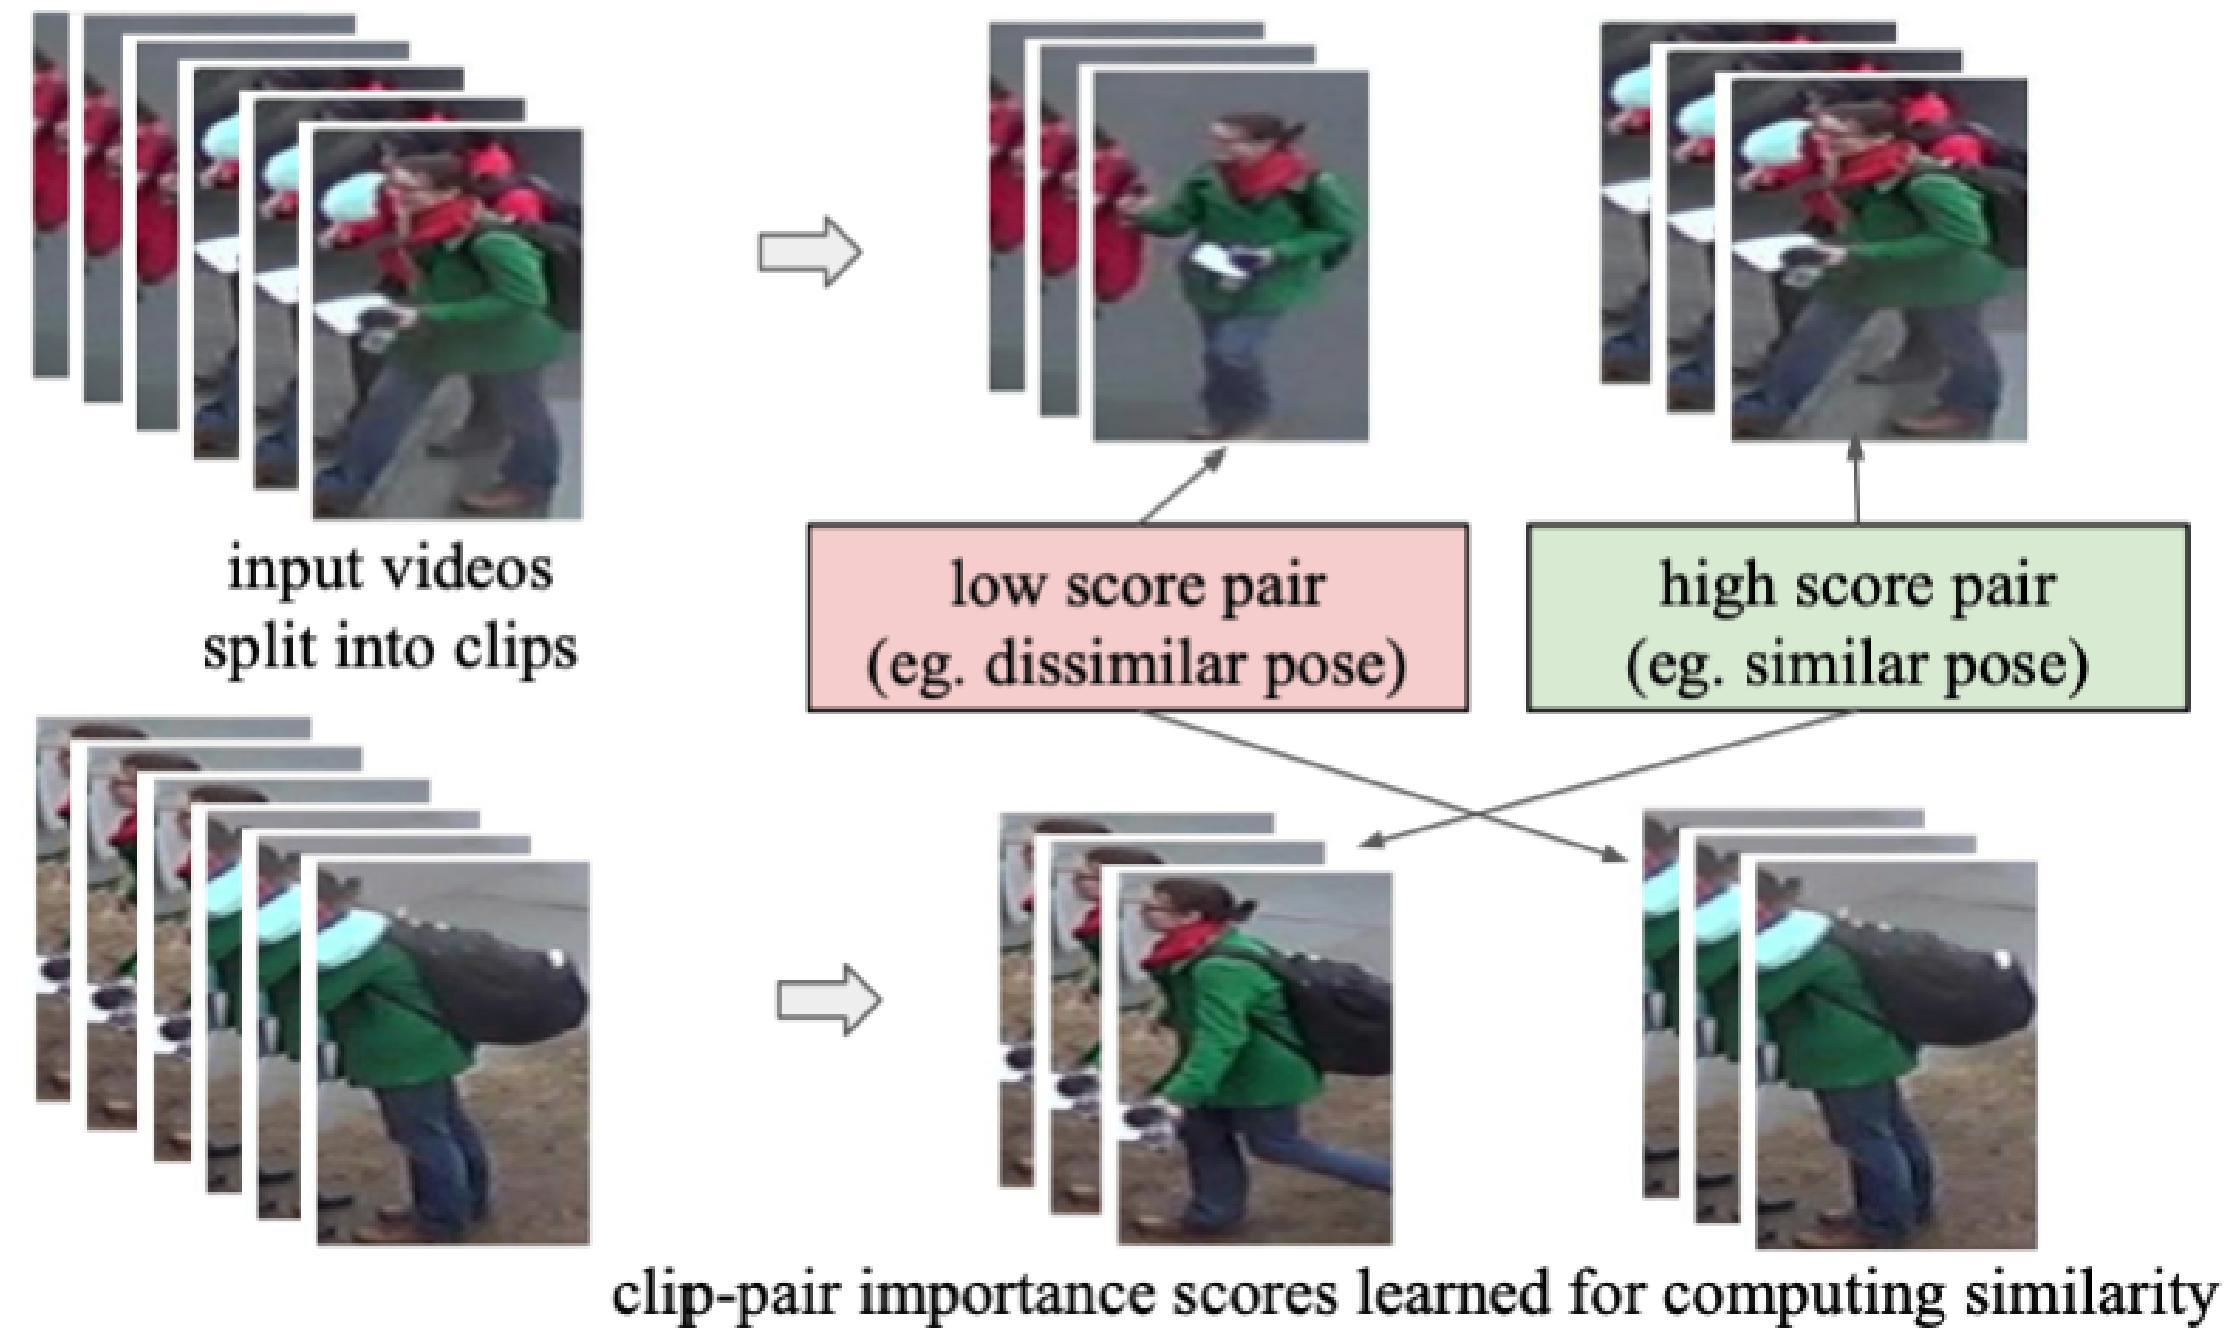
1. **3D CNN:** Allows the network to learn spatial and temporal features jointly (Clip-level instead of Frame-level)
2. **Split video into short clips:** The video is too long.
3. **Pair each clip:** $M \times M \rightarrow M$ is a number of clips.
4. **Compute importance for each pair:** Using MLP network. This score indicates how informative that clip pair is for estimating video similarity

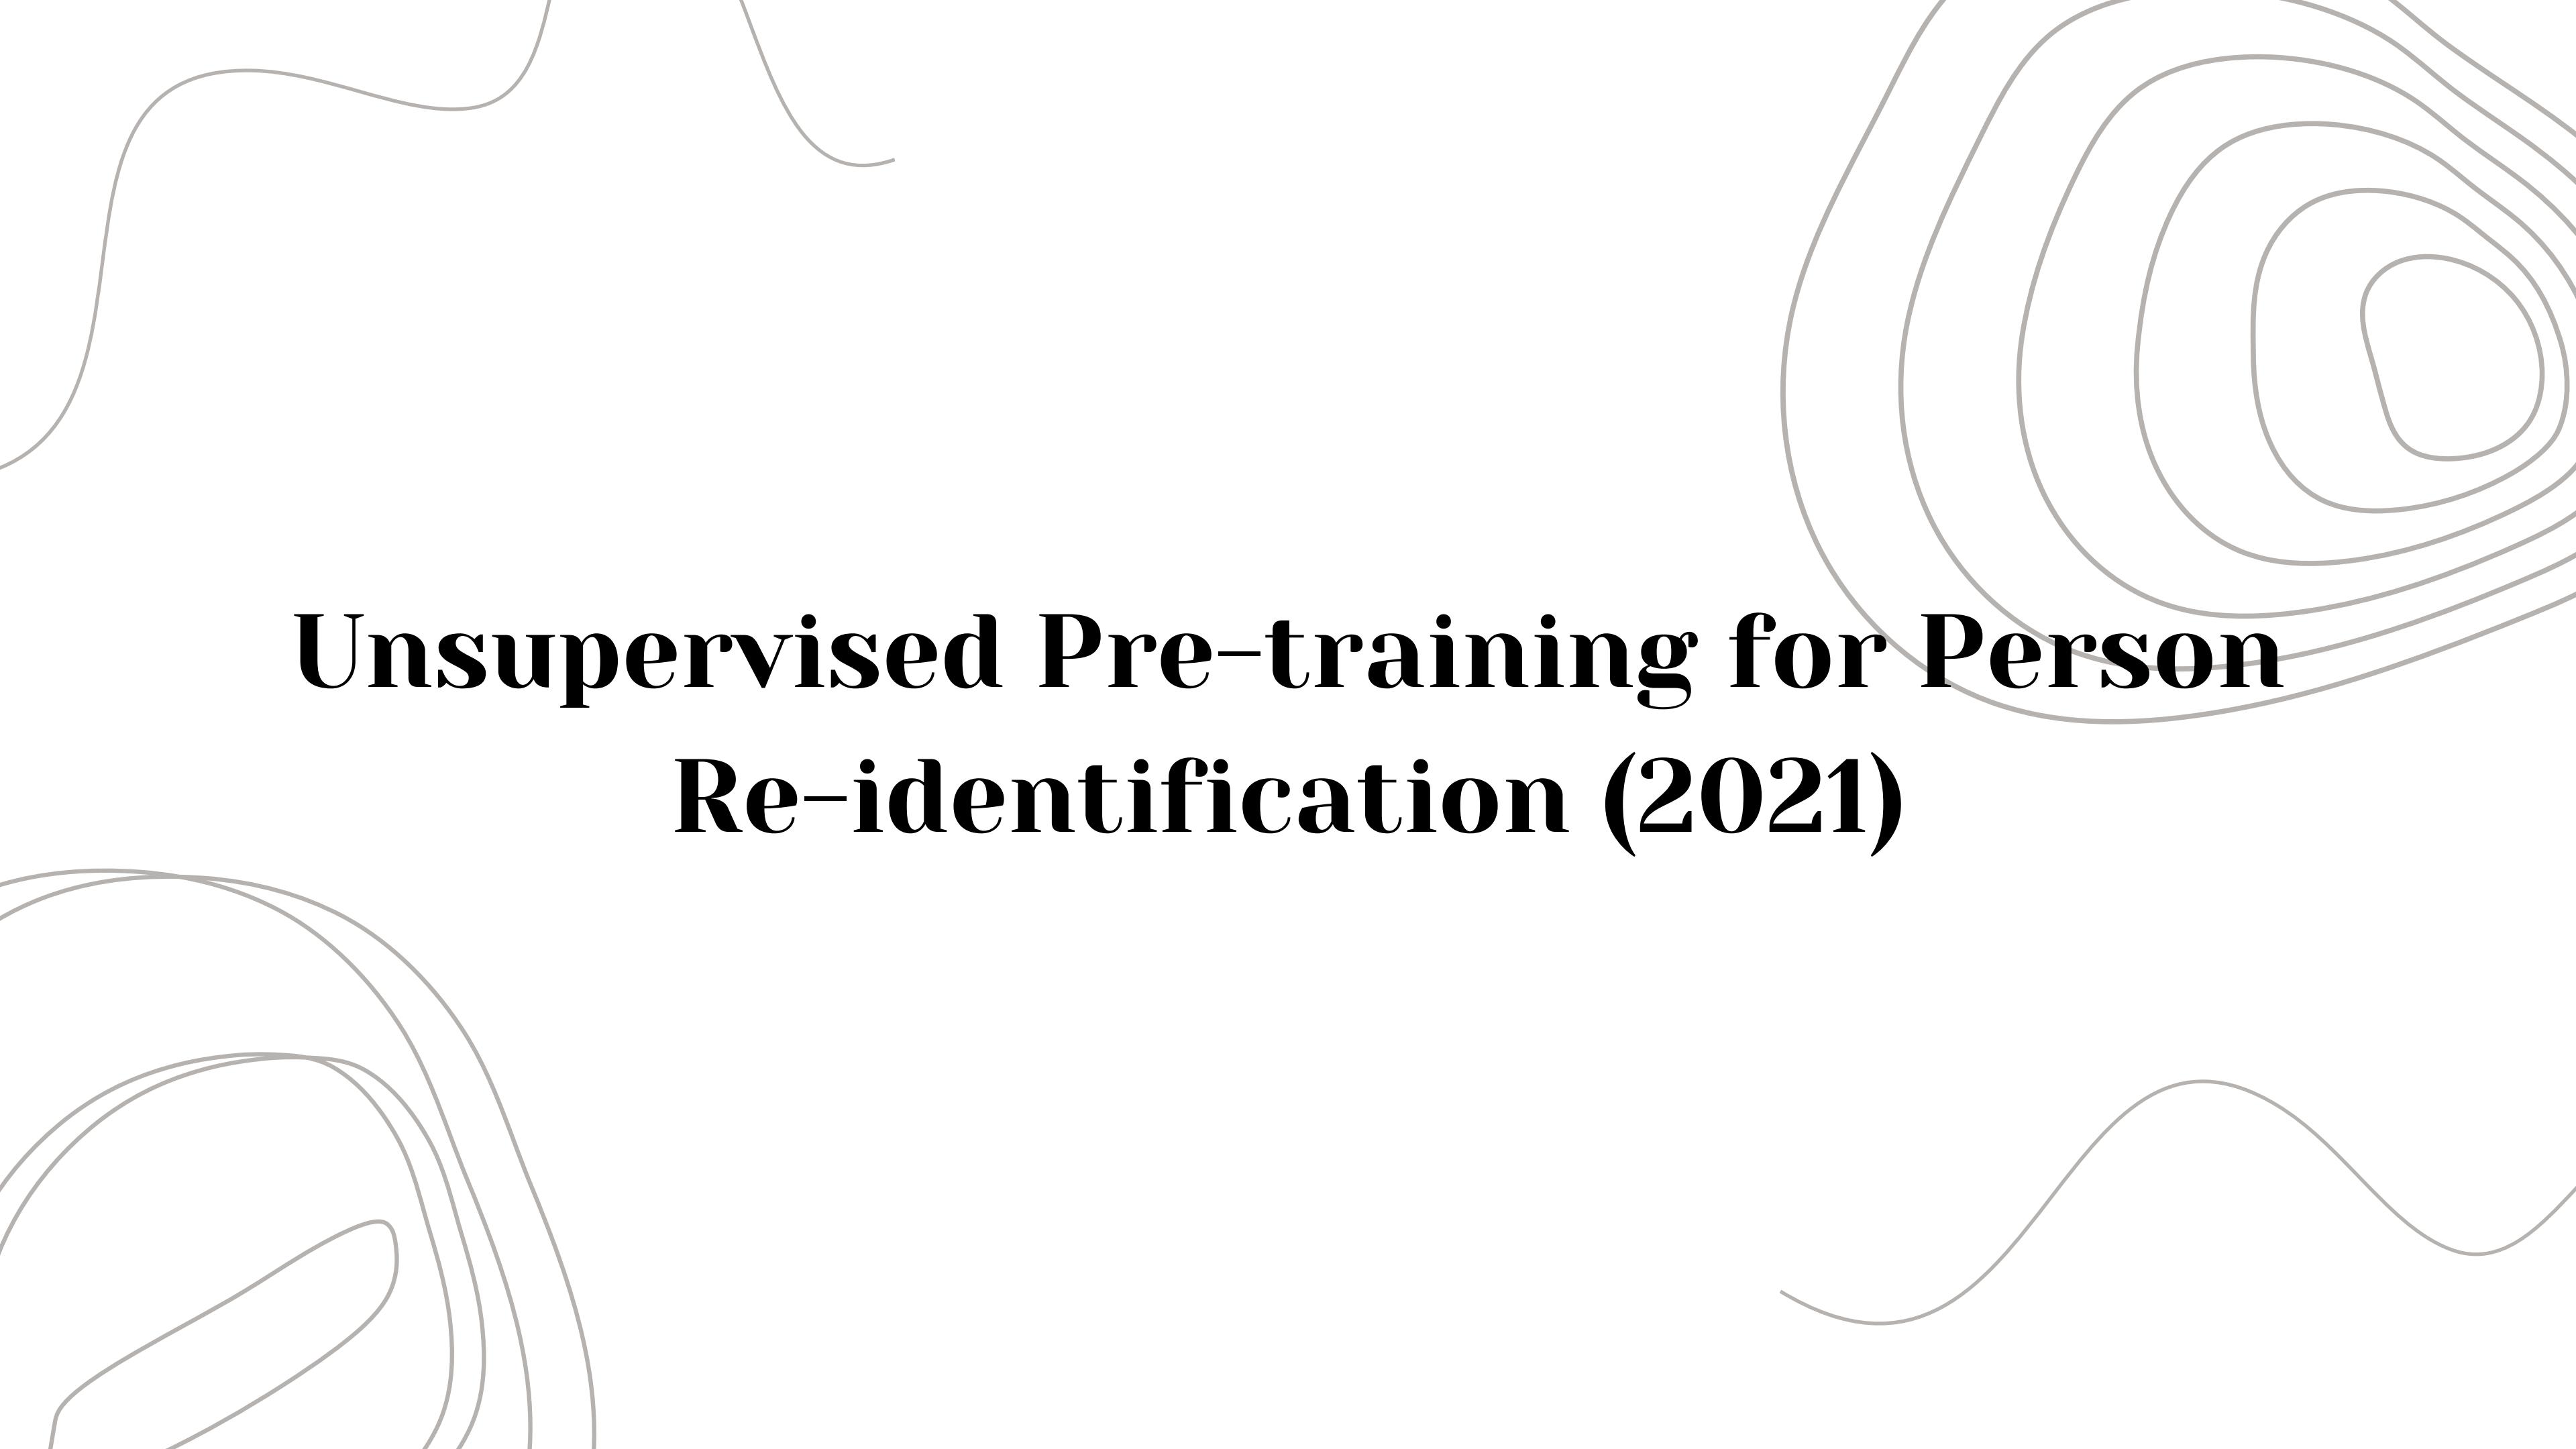
Steps

5. **Train the model:** The model learns to focus on the most informative clip pairs for re-identification, while automatically downweighting the contribution of noisy or uninformative clips.

(Weighted Average)

Steps





Unsupervised Pre-training for Person Re-identification (2021)

Steps

1. **LUPerson:** The authors build a large-scale unlabeled person Re-ID dataset called "LUPerson" containing over 4 million images of 200K identities.
2. **MoCoV2:** Using contrastive learning (One of the self-supervised learning)
3. **Direct applying MoCoV2 doesn't work:** Data augmentation, tuning the model parameters.

ByteTrack

ByteTrack

One of the algorithms for tracking the object in the video.

How does ByteTrack work?

- **Object detection:** Use YoloX object detection model as its backbone (YoloX is a single-stage object detector)
- **Feature Extraction**
- **Compute association score:** Computes an association score between each detected object in the current frame and the tracked objects from the previous frames by considers the spatial proximity (IoU) and appearance similarity (cosine distance) between the objects.

How does ByteTrack work?

- **A hierarchical association:** Handle occlusions and missing detections.
- **Track management:**
 - Maintain active track (Unique object)
 - When detect new object -> Assign a new track
 - When object disappear for a certain time ->
 - Remove from a track

How does ByteTrack work?

- **Motion Model**
 - Employs a simple motion model to predict the position of tracked objects in the next frame.
 - The motion model assumes a constant velocity and updates the object's position based on its previous location and velocity.

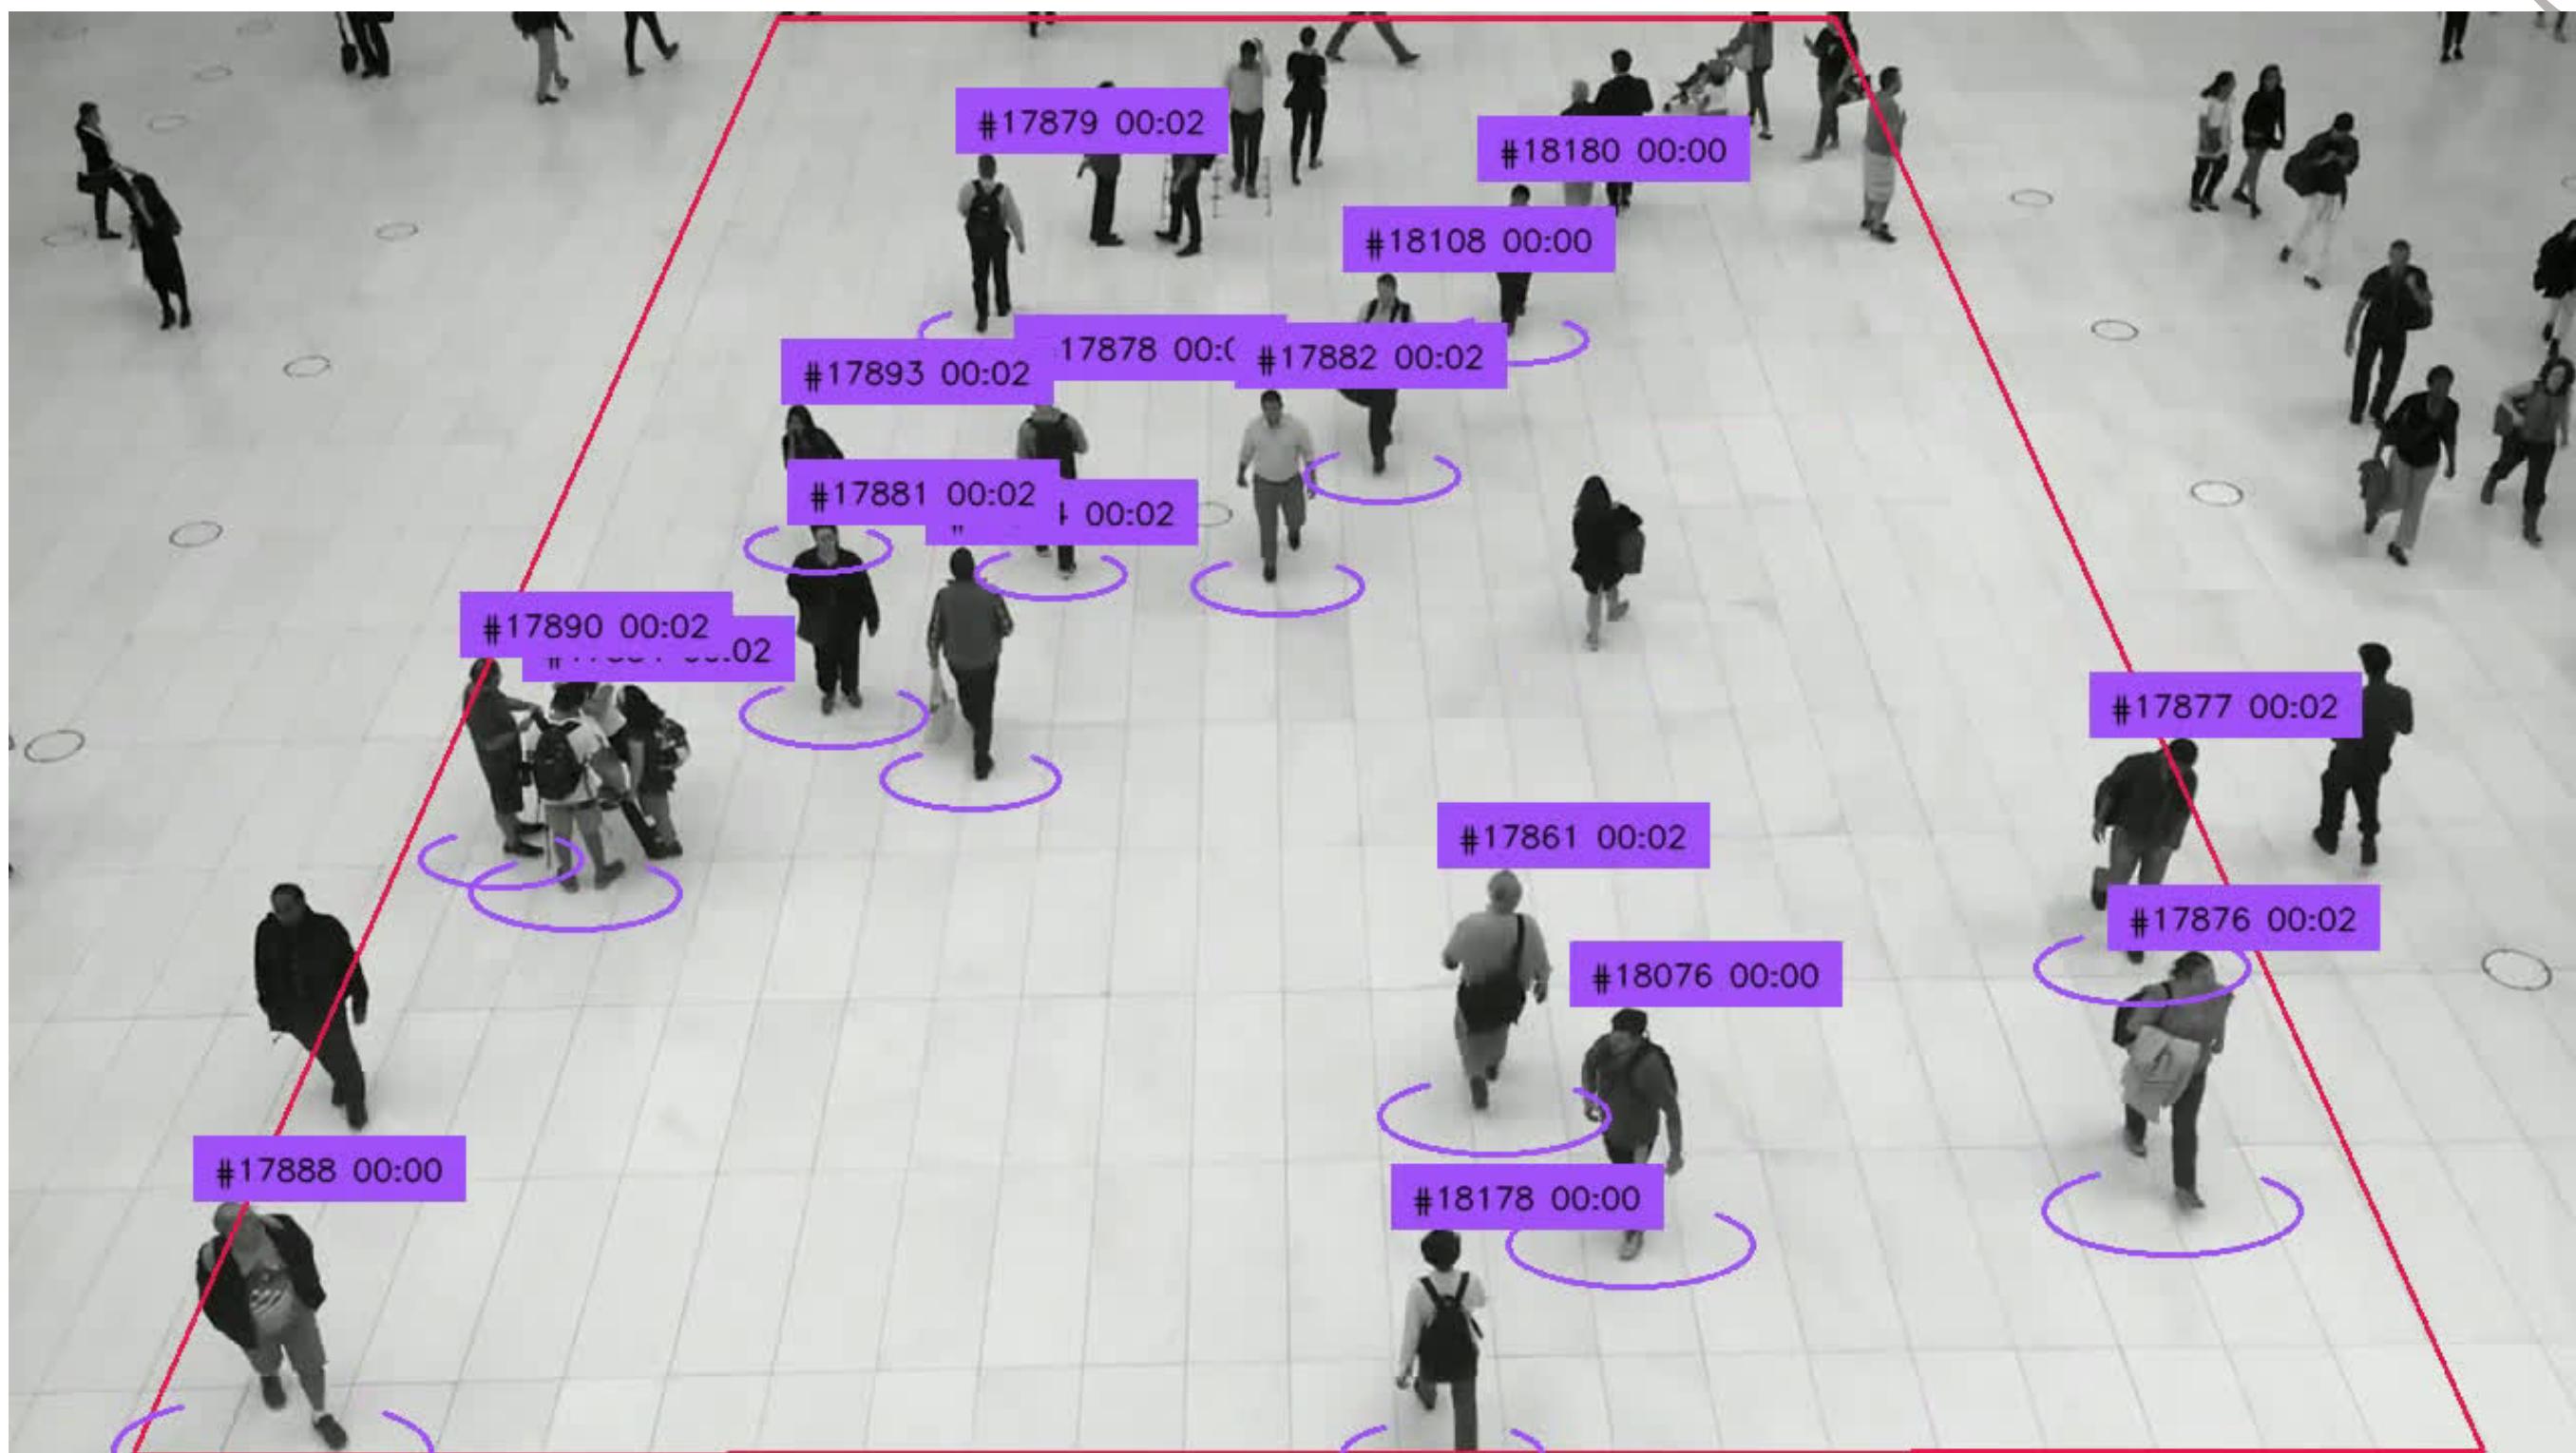
Implementation



Implementation



Implementation



Active Learning

Active Learning



Create a project

roboflow

Let's create your project.

FashionTest > [New Public Project](#)

Project Name: Fashion

License: CC BY 4.0

Annotation Group: fashions

Project Type:

- Object Detection**
Identify objects and their positions with bounding boxes.
Best For: # Counting, ⚡ Tracking
- Classification**
Assign labels to the entire image.
Classification Type: Multi-Label (selected), Single-Label
Best For: Filtering, ✓ Content Moderation
- Instance Segmentation**
Detect multiple objects and their actual shape.
Best For: Measurements, ⚡ Odd Shapes
- Keypoint Detection**
Identify keypoints ("skeletons") to subjects.
Best For: ⚡ Pose Estimation

Show More ↓

Cancel [Create Public Project](#)

Upload to Roboflow

```
dir_name = f"{HOME}/yolov9/mytest"
file_extension_type = [".jpg", ".png", ".jpeg", ".webp"]

image_glob = [file for ext_type in file_extension_type for file in glob.glob(dir_name + '/*' + ext_type)]

for image_path in image_glob:
    annot_path = image_path.rsplit(".", 1)[0] + ".txt"
    annot_path = annot_path if os.path.exists(annot_path) else None
    print(project.single_upload(
        image_path=image_path,
        annotation_path=annot_path,
    ))

{'image': {'success': True, 'id': 'u1WIOKPk1gN7YBcDVloF'}, 'annotation': {'success': True}, 'upload_time': 5.829010963439941, 'annotation_time': 1.0019919872283936}
{'image': {'success': True, 'id': 'drzteocMALyStRLbKZcC'}, 'annotation': {'success': True}, 'upload_time': 4.348002672195435, 'annotation_time': 1.123000144958496}
{'image': {'success': True, 'id': 'HOXR5hYMFx48V4W7Xzt6'}, 'annotation': {'success': True}, 'upload_time': 2.7090094089508057, 'annotation_time': 0.8689875602722168}
{'image': {'success': True, 'id': 'qRpqjSG9zBcsLgBQDMJ'}, 'annotation': None, 'upload_time': 2.6800103187561035, 'annotation_time': None}
{'image': {'success': True, 'id': 'HIBtCUxRyPYJ0406cXTt'}, 'annotation': {'success': True}, 'upload_time': 0.9860019683837891, 'annotation_time': 0.9089875221252441}
{'image': {'success': True, 'id': 'XFoOtp111UmgWsBkit6E'}, 'annotation': {'success': True}, 'upload_time': 1.8720018863677979, 'annotation_time': 0.8840103149414062}
{'image': {'success': True, 'id': 'KZpD2s4fj0aAJYlja10U'}, 'annotation': {'success': True, 'warnings': True, 'badAnnotations': [{"type': 'warning', 'code': 'annotation:trimmed', 'key': 'Screenshot_20240408_123727_Gallery.jpg', 'label': '10'}, {"type': 'warning', 'code': 'annotation:trimmed', 'key': 'Screenshot_20240408_123727_Gallery.jpg', 'label': '10'}]}, 'upload_time': 2.6599998474121094, 'annotation_time': 0.8830046653747559}
{'image': {'success': True, 'id': 'HCStD9fOMLBd1CKutsde'}, 'annotation': {'success': True}, 'upload_time': 2.871997356414795, 'annotation_time': 0.8659868240356445}
{'image': {'success': True, 'id': 'XxqxLdmT8P2BeV70zARY'}, 'annotation': {'success': True}, 'upload_time': 3.119680643081665, 'annotation_time': 0.9110023975372314}
{'image': {'success': True, 'id': 'CEUEBttQAOgQyLgKqwYI'}, 'annotation': {'success': True}, 'upload_time': 2.4799888134002686, 'annotation_time': 1.313999891281128}
```



Demo